

Personalized cancer diagnosis

Exploratory Data Analysis

In [2]:

```
import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import numpy as np
from nltk.corpus import stopwords
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.manifold import TSNE
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics.classification import accuracy_score, log_loss
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import SGDClassifier
from imblearn.over_sampling import SMOTE
from collections import Counter
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
import math
from sklearn.metrics import normalized_mutual_info_score
from sklearn.ensemble import RandomForestClassifier
warnings.filterwarnings("ignore")

from mlxtend.classifier import StackingClassifier

from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
```

3.1. Reading Data

3.1.1. Reading Gene and Variation Data

In [3]:

```
data = pd.read_csv('training/training_variants')
print('Number of data points : ', data.shape[0])
print('Number of features : ', data.shape[1])
print('Features : ', data.columns.values)
data.head()
```

```
Number of data points : 3321
Number of features : 4
Features : ['ID' 'Gene' 'Variation' 'Class']
```

Out[3]:

	ID	Gene	Variation	Class
0	0	FAM58A	Truncating Mutations	1

1	ID	Gene	Variation	Class
2	2	CBL	Q249E	2
3	3	CBL	N454D	3
4	4	CBL	L399V	4

training/training_variants is a comma separated file containing the description of the genetic mutations used for training. Fields are

- **ID** : the id of the row used to link the mutation to the clinical evidence
- **Gene** : the gene where this genetic mutation is located
- **Variation** : the aminoacid change for this mutations
- **Class** : 1-9 the class this genetic mutation has been classified on

3.1.2. Reading Text Data

In [4]:

```
# note the separator in this file
data_text = pd.read_csv("training/training_text", sep="\\|\\|", engine="python", names=["ID", "TEXT"], skip
rows=1)
print('Number of data points : ', data_text.shape[0])
print('Number of features : ', data_text.shape[1])
print('Features : ', data_text.columns.values)
data_text.head()
```

```
Number of data points : 3321
Number of features : 2
Features : ['ID' 'TEXT']
```

Out[4]:

	ID	TEXT
0	0	Cyclin-dependent kinases (CDKs) regulate a var...
1	1	Abstract Background Non-small cell lung canc...
2	2	Abstract Background Non-small cell lung canc...
3	3	Recent evidence has demonstrated that acquired...
4	4	Oncogenic mutations in the monomeric Casitas B...

3.1.3. Preprocessing of text

In [5]:

```
# loading stop words from nltk library
stop_words = set(stopwords.words('english'))

def nlp_preprocessing(total_text, index, column):
    if type(total_text) is not int:
        string = ""
        # replace every special char with space
        total_text = re.sub('[^a-zA-Z0-9\\n]', ' ', total_text)
        # replace multiple spaces with single space
        total_text = re.sub('\\s+', ' ', total_text)
        # converting all the chars into lower-case.
        total_text = total_text.lower()

        for word in total_text.split():
            # if the word is a not a stop word then retain that word from the data
            if not word in stop_words:
                string += word + " "
```

```
data_text[column][index] = string
```

In [6]:

```
#text processing stage.
start_time = time.clock()
for index, row in data_text.iterrows():
    if type(row['TEXT']) is str:
        nlp_preprocessing(row['TEXT'], index, 'TEXT')
    else:
        print("there is no text description for id:",index)
print('Time took for preprocessing the text :',time.clock() - start_time, "seconds")
```

```
there is no text description for id: 1109
there is no text description for id: 1277
there is no text description for id: 1407
there is no text description for id: 1639
there is no text description for id: 2755
Time took for preprocessing the text : 194.607563322 seconds
```

In [7]:

```
#merging both gene_variations and text data based on ID
result = pd.merge(data, data_text,on='ID', how='left')
result.head()
```

Out[7]:

	ID	Gene	Variation	Class	TEXT
0	0	FAM58A	Truncating Mutations	1	cyclin dependent kinases cdks regulate variety...
1	1	CBL	W802*	2	abstract background non small cell lung cancer...
2	2	CBL	Q249E	2	abstract background non small cell lung cancer...
3	3	CBL	N454D	3	recent evidence demonstrated acquired uniparen...
4	4	CBL	L399V	4	oncogenic mutations monomeric casitas b lineag...

In [8]:

```
result[result.isnull().any(axis=1)]
```

Out[8]:

	ID	Gene	Variation	Class	TEXT
1109	1109	FANCA	S1088F	1	NaN
1277	1277	ARID5B	Truncating Mutations	1	NaN
1407	1407	FGFR3	K508M	6	NaN
1639	1639	FLT1	Amplification	6	NaN
2755	2755	BRAF	G596C	7	NaN

In [9]:

```
result.loc[result['TEXT'].isnull(), 'TEXT'] = result['Gene'] + ' ' + result['Variation']
```

In [10]:

```
result[result['ID']==1109]
```

Out[10]:

	ID	Gene	Variation	Class	TEXT
1109	1109	FANCA	S1088F	1	FANCA S1088F

3.1.4. Test, Train and Cross Validation Split

3.1.4.1. Splitting data into train, test and cross validation (64:20:16)

In [11]:

```
y_true = result['Class'].values
result.Gene = result.Gene.str.replace('\s+', '_')
result.Variation = result.Variation.str.replace('\s+', '_')

# split the data into test and train by maintaining same distribution of output variable 'y_true'
[stratify=y_true]
X_train, test_df, y_train, y_test = train_test_split(result, y_true, stratify=y_true, test_size=0.2)
# split the train data into train and cross validation by maintaining same distribution of output
variable 'y_train' [stratify=y_train]
train_df, cv_df, y_train, y_cv = train_test_split(X_train, y_train, stratify=y_train, test_size=0.2)
```

We split the data into train, test and cross validation data sets, preserving the ratio of class distribution in the original data set

In [12]:

```
print('Number of data points in train data:', train_df.shape[0])
print('Number of data points in test data:', test_df.shape[0])
print('Number of data points in cross validation data:', cv_df.shape[0])
```

```
Number of data points in train data: 2124
Number of data points in test data: 665
Number of data points in cross validation data: 532
```

3.1.4.2. Distribution of y_i's in Train, Test and Cross Validation datasets

In [13]:

```
# it returns a dict, keys as class labels and values as the number of data points in that class
train_class_distribution = train_df['Class'].value_counts().sort_index()
test_class_distribution = test_df['Class'].value_counts().sort_index()
cv_class_distribution = cv_df['Class'].value_counts().sort_index()

my_colors = 'rgbkymc'
train_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of y_i in train data')
plt.grid()
plt.show()

# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html
# -(train_class_distribution.values): the minus sign will give us in decreasing order
sorted_yi = np.argsort(-train_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ':', train_class_distribution.values[i], '(', np.round((train_class_distribution.values[i]/train_df.shape[0]*100), 3), '%)')

print('-'*80)
my_colors = 'rgbkymc'
test_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of y_i in test data')
plt.grid()
plt.show()

# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html
# -(train_class_distribution.values): the minus sign will give us in decreasing order
sorted_yi = np.argsort(-test_class_distribution.values)
```


```

for i in sorted_yi:
    print('Number of data points in class', i+1, ':', test_class_distribution.values[i], '(', np.round(
nd((test_class_distribution.values[i]/test_df.shape[0]*100), 3), '%)')


print('-'*80)
my_colors = 'rgbkymc'
cv_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of yi in cross validation data')
plt.grid()
plt.show()

# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html
# -(train_class_distribution.values): the minus sign will give us in decreasing order
sorted_yi = np.argsort(-train_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ':', cv_class_distribution.values[i], '(', np.round(
((cv_class_distribution.values[i]/cv_df.shape[0]*100), 3), '%)')


```



Number of data points in class 7 : 609 (28.672 %)
 Number of data points in class 4 : 439 (20.669 %)
 Number of data points in class 1 : 363 (17.09 %)
 Number of data points in class 2 : 289 (13.606 %)
 Number of data points in class 6 : 176 (8.286 %)
 Number of data points in class 5 : 155 (7.298 %)
 Number of data points in class 3 : 57 (2.684 %)
 Number of data points in class 9 : 24 (1.13 %)
 Number of data points in class 8 : 12 (0.565 %)



Number of data points in class 7 : 191 (28.722 %)
 Number of data points in class 4 : 137 (20.602 %)
 Number of data points in class 1 : 114 (17.143 %)
 Number of data points in class 2 : 91 (13.684 %)
 Number of data points in class 6 : 55 (8.271 %)
 Number of data points in class 5 : 48 (7.218 %)
 Number of data points in class 3 : 18 (2.707 %)
 Number of data points in class 9 : 7 (1.053 %)
 Number of data points in class 8 : 4 (0.602 %)



Number of data points in class 7 : 153 (28.759 %)
 Number of data points in class 4 : 110 (20.677 %)
 Number of data points in class 1 : 91 (17.105 %)
 Number of data points in class 2 : 72 (13.534 %)
 Number of data points in class 6 : 44 (8.271 %)
 Number of data points in class 5 : 39 (7.331 %)
 Number of data points in class 3 : 14 (2.632 %)
 Number of data points in class 9 : 6 (1.128 %)
 Number of data points in class 8 : 3 (0.564 %)

3.2 Prediction using a 'Random' Model

In a 'Random' Model, we generate the NINE class probabilities randomly such that they sum to 1.

In [14]:

```

# This function plots the confusion matrices given y_i, y_i_hat.
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)
    # C = 9,9 matrix, each cell (i,j) represents number of points of class i are predicted class j

    A = ((C.T)/(C.sum(axis=1))).T
    #divide each element of the confusion matrix with the sum of elements in that column

    # C = [[1, 2],
    #       [3, 4]]
    # C.T = [[1, 3],
    #         [2, 4]]
    # C.sum(axis = 1)  axis=0 corresponds to columns and axis=1 corresponds to rows in two
    # dimensional array
    # C.sum(axix =1) = [[3, 7]]
    # ((C.T)/(C.sum(axis=1))) = [[1/3, 3/7]
    #                             [2/3, 4/7]]

    # ((C.T)/(C.sum(axis=1))).T = [[1/3, 2/3]
    #                               [3/7, 4/7]]
    # sum of row elements = 1

    B = (C/C.sum(axis=0))
    #divide each element of the confusion matrix with the sum of elements in that row
    # C = [[1, 2],
    #       [3, 4]]
    # C.sum(axis = 0)  axis=0 corresponds to columns and axis=1 corresponds to rows in two
    # dimensional array
    # C.sum(axix =0) = [[4, 6]]
    # (C/C.sum(axis=0)) = [[1/4, 2/6],
    #                       [3/4, 4/6]]

    labels = [1,2,3,4,5,6,7,8,9]
    # representing A in heatmap format
    print("-"*20, "Confusion matrix", "-"*20)
    plt.figure(figsize=(20,7))
    sns.heatmap(C, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
  
```

```

plt.ylabel('Original Class')
plt.show()

print("-"*20, "Precision matrix (Column Sum=1)", "-"*20)
plt.figure(figsize=(20,7))
sns.heatmap(B, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

# representing B in heatmap format
print("-"*20, "Recall matrix (Row sum=1)", "-"*20)
plt.figure(figsize=(20,7))
sns.heatmap(A, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

```

In [15]:

```

# we need to generate 9 numbers and the sum of numbers should be 1
# one solution is to generate 9 numbers and divide each of the numbers by their sum
# ref: https://stackoverflow.com/a/18662466/4084039
test_data_len = test_df.shape[0]
cv_data_len = cv_df.shape[0]

# we create a output array that has exactly same size as the CV data
cv_predicted_y = np.zeros((cv_data_len,9))
for i in range(cv_data_len):
    rand_probs = np.random.rand(1,9)
    cv_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Cross Validation Data using Random Model",log_loss(y_cv,cv_predicted_y, eps=1e-15))

# Test-Set error.
#we create a output array that has exactly same as the test data
test_predicted_y = np.zeros((test_data_len,9))
for i in range(test_data_len):
    rand_probs = np.random.rand(1,9)
    test_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Test Data using Random Model",log_loss(y_test,test_predicted_y, eps=1e-15))


predicted_y = np.argmax(test_predicted_y, axis=1)
plot_confusion_matrix(y_test, predicted_y+1)

```


Log loss on Cross Validation Data using Random Model 2.4851797771943605

Log loss on Test Data using Random Model 2.4689181852544024


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



3.3 Univariate Analysis

In [16]:

```
# code for response coding with Laplace smoothing.
# alpha : used for laplace smoothing
# feature: ['gene', 'variation']
# df: ['train_df', 'test_df', 'cv_df']
# algorithm
# -----
# Consider all unique values and the number of occurrences of given feature in train data dataframe
# build a vector (1*9) , the first element = (number of times it occurred in class1 + 10*alpha / number of times it occurred in total data+90*alpha)
# gv_dict is like a look up table, for every gene it store a (1*9) representation of it
# for a value of feature in df:
# if it is in train data:
# we add the vector that was stored in 'gv_dict' look up table to 'gv_fea'
# if it is not there is train:
# we add [1/9, 1/9, 1/9, 1/9, 1/9, 1/9, 1/9, 1/9, 1/9] to 'gv_fea'
# return 'gv_fea'
# -----

# get_gv_fea_dict: Get Gene variation Feature Dict
def get_gv_fea_dict(alpha, feature, df):
    # value_count: it contains a dict like
    # print(train_df['Gene'].value_counts())
    # output:
    #          {BRCA1      174
    #          TP53       106
    #          EGFR        86
```



```

#         EGFR           86
#         BRCA2          75
#         PTEN           69
#         KIT            61
#         BRAF           60
#         ERBB2          47
#         PDGFRA         46
#         ...}
# print(train_df['Variation'].value_counts())
# output:
# {
# Truncating_Mutations           63
# Deletion                       43
# Amplification                  43
# Fusions                        22
# Overexpression                  3
# E17K                           3
# Q61L                           3
# S222D                           2
# P130S                           2
# ...
# }
value_count = train_df[feature].value_counts()

# gv_dict : Gene Variation Dict, which contains the probability array for each gene/variation
gv_dict = dict()

# denominator will contain the number of time that particular feature occurred in whole data
for i, denominator in value_count.items():
    # vec will contain (p(yi==1/Gi) probability of gene/variation belongs to particular class
    # vec is 9 dimensional vector
    vec = []
    for k in range(1,10):
        # print(train_df.loc[(train_df['Class']==1) & (train_df['Gene']=='BRCA1')])
        #
        # ID      Gene      Variation      Class
        # 2470    2470    BRCA1      S1715C      1
        # 2486    2486    BRCA1      S1841R      1
        # 2614    2614    BRCA1      M1R        1
        # 2432    2432    BRCA1      L1657P      1
        # 2567    2567    BRCA1      T1685A      1
        # 2583    2583    BRCA1      E1660G      1
        # 2634    2634    BRCA1      W1718L      1
        # cls_cnt.shape[0] will return the number of rows

        cls_cnt = train_df.loc[(train_df['Class']==k) & (train_df[feature]==i)]

        # cls_cnt.shape[0] (numerator) will contain the number of time that particular feature occurred in whole data
        vec.append((cls_cnt.shape[0] + alpha*10)/ (denominator + 90*alpha))

    # we are adding the gene/variation to the dict as key and vec as value
    gv_dict[i]=vec
return gv_dict

# Get Gene variation feature
def get_gv_feature(alpha, feature, df):
    # print(gv_dict)
    #
    # {'BRCA1': [0.20075757575757575, 0.03787878787878788, 0.06818181818181817,
0.13636363636363635, 0.25, 0.19318181818181818, 0.03787878787878788, 0.03787878787878788,
0.03787878787878788],
#
# 'TP53': [0.32142857142857145, 0.061224489795918366, 0.061224489795918366,
0.27040816326530615, 0.061224489795918366, 0.066326530612244902, 0.051020408163265307, 0.051020408
163265307, 0.056122448979591837],
#
# 'EGFR': [0.056818181818181816, 0.21590909090909091, 0.0625, 0.06818181818181817,
0.068181818181817, 0.0625, 0.34659090909090912, 0.0625, 0.0568181818181816],
#
# 'BRCA2': [0.13333333333333333, 0.060606060606060608, 0.060606060606060608,
0.07878787878787878, 0.13939393939393939, 0.34545454545454546, 0.060606060606060608,
0.060606060606060608, 0.060606060606060608],
#
# 'PTEN': [0.069182389937106917, 0.062893081761006289, 0.069182389937106917,
0.46540880503144655, 0.075471698113207544, 0.062893081761006289, 0.069182389937106917, 0.062893081
761006289, 0.062893081761006289],
#
# 'KIT': [0.066225165562913912, 0.25165562913907286, 0.072847682119205295,
0.072847682119205295, 0.066225165562913912, 0.066225165562913912, 0.27152317880794702,
0.066225165562913912, 0.066225165562913912],
#
# 'BRAF': [0.066666666666666666, 0.17999999999999999, 0.073333333333333334,
0.073333333333333334, 0.093333333333333338, 0.080000000000000002, 0.29999999999999999,
0.066666666666666666, 0.066666666666666666],

```

```

#     ...
#     }
gv_dict = get_gv_fea_dict(alpha, feature, df)
# value_count is similar in get_gv_fea_dict
value_count = train_df[feature].value_counts()

# gv_fea: Gene_variation feature, it will contain the feature for each feature value in the data
gv_fea = []
# for every feature values in the given data frame we will check if it is there in the train data then we will add the feature to gv_fea
# if not we will add [1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9] to gv_fea
for index, row in df.iterrows():
    if row[feature] in dict(value_count).keys():
        gv_fea.append(gv_dict[row[feature]])
    else:
        gv_fea.append([1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9])
#     gv_fea.append([-1,-1,-1,-1,-1,-1,-1,-1,-1,-1])
return gv_fea

```

when we calculate the probability of a feature belongs to any particular class, we apply laplace smoothing

- $(\text{numerator} + 10 \cdot \alpha) / (\text{denominator} + 90 \cdot \alpha)$

3.2.1 Univariate Analysis on Gene Feature

Q1. Gene, What type of feature it is ?

Ans. Gene is a categorical variable

Q2. How many categories are there and How they are distributed?

In [17]:

```

unique_genes = train_df['Gene'].value_counts()
print('Number of Unique Genes :', unique_genes.shape[0])
# the top 10 genes that occurred most
print(unique_genes.head(10))

```

```

Number of Unique Genes : 230
BRCA1    159
TP53     97
EGFR     97
BRCA2    90
PTEN     79
KIT       64
BRAF     58
ERBB2    49
ALK       42
TSC2     38
Name: Gene, dtype: int64

```

In [18]:

```

print("Ans: There are", unique_genes.shape[0], "different categories of genes in the train data, and they are distributed as follows",)

```


Ans: There are 230 different categories of genes in the train data, and they are distributed as follows

In [19]:

```


s = sum(unique_genes.values);
h = unique_genes.values/s;
plt.plot(h, label="Histogram of Genes")
plt.xlabel('Index of a Gene')
plt.ylabel('Number of Occurrences')
plt.legend()
plt.grid()
plt.show()

```



In [20]:

```
c = np.cumsum(h)
plt.plot(c, label='Cumulative distribution of Genes')
plt.grid()
plt.legend()
plt.show()
```



Q3. How to featurize this Gene feature ?

Ans. there are two ways we can featurize this variable check out this video:

<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>

1. One hot Encoding
2. Response coding

We will choose the appropriate featurization based on the ML model we use. For this problem of multi-class classification with categorical features, one-hot encoding is better for Logistic regression while response coding is better for Random Forests.

In [21]:

```
#response-coding of the Gene feature
# alpha is used for laplace smoothing
alpha = 1
# train gene feature
train_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", train_df))
# test gene feature
test_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", test_df))
# cross validation gene feature
cv_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", cv_df))
```

In [22]:

```
print("train_gene_feature_responseCoding is converted feature using response coding method. The shape of gene feature:", train_gene_feature_responseCoding.shape)
```

train_gene_feature_responseCoding is converted feature using response coding method. The shape of g

```
ene feature: (2124, 9)
```

In [23]:

```
# one-hot encoding of Gene feature.
gene_vectorizer = CountVectorizer()
train_gene_feature_onehotCoding = gene_vectorizer.fit_transform(train_df['Gene'])
test_gene_feature_onehotCoding = gene_vectorizer.transform(test_df['Gene'])
cv_gene_feature_onehotCoding = gene_vectorizer.transform(cv_df['Gene'])
```

In [24]:

```
train_df['Gene'].head()
```

Out[24]:

```
1384    FGFR1
2808    BRCA2
307     H3F3A
1933      SMO
208     EGFR
Name: Gene, dtype: object
```

In [25]:

```
gene_vectorizer.get_feature_names()
```

Out[25]:

```
['abl1',
 'acvr1',
 'ago2',
 'akt1',
 'akt2',
 'akt3',
 'alk',
 'apc',
 'ar',
 'araf',
 'arid1a',
 'arid1b',
 'arid2',
 'arid5b',
 'asxl2',
 'atm',
 'atrx',
 'aurka',
 'aurkb',
 'axin1',
 'b2m',
 'bap1',
 'bcl10',
 'bcl2l1',
 'bcor',
 'braf',
 'brca1',
 'brca2',
 'brd4',
 'brip1',
 'btk',
 'card11',
 'carm1',
 'casp8',
 'cbl',
 'ccnd1',
 'ccnd2',
 'ccnd3',
 'ccne1',
 'cdh1',
 'cdk12',
 'cdk4',
 'cdk6',
 'cdkn1a',
 'cdkn1b',
 'cdkn1c',
 'cdkn1d',
 'cdkn1e',
 'cdkn1f',
 'cdkn1g',
 'cdkn1h',
 'cdkn1i',
 'cdkn1j',
 'cdkn1k',
 'cdkn1l',
 'cdkn1m',
 'cdkn1n',
 'cdkn1o',
 'cdkn1p',
 'cdkn1q',
 'cdkn1r',
 'cdkn1s',
 'cdkn1t',
 'cdkn1u',
 'cdkn1v',
 'cdkn1w',
 'cdkn1x',
 'cdkn1y',
 'cdkn1z',
 'cdkn1aa',
 'cdkn1ab',
 'cdkn1ac',
 'cdkn1ad',
 'cdkn1ae',
 'cdkn1af',
 'cdkn1ag',
 'cdkn1ah',
 'cdkn1ai',
 'cdkn1aj',
 'cdkn1ak',
 'cdkn1al',
 'cdkn1am',
 'cdkn1an',
 'cdkn1ao',
 'cdkn1ap',
 'cdkn1aq',
 'cdkn1ar',
 'cdkn1as',
 'cdkn1at',
 'cdkn1au',
 'cdkn1av',
 'cdkn1aw',
 'cdkn1ax',
 'cdkn1ay',
 'cdkn1az',
 'cdkn1ba',
 'cdkn1bb',
 'cdkn1bc',
 'cdkn1bd',
 'cdkn1be',
 'cdkn1bf',
 'cdkn1bg',
 'cdkn1bh',
 'cdkn1bi',
 'cdkn1bj',
 'cdkn1bk',
 'cdkn1bl',
 'cdkn1bm',
 'cdkn1bn',
 'cdkn1bo',
 'cdkn1bp',
 'cdkn1bq',
 'cdkn1br',
 'cdkn1bs',
 'cdkn1bt',
 'cdkn1bu',
 'cdkn1bv',
 'cdkn1bw',
 'cdkn1bx',
 'cdkn1by',
 'cdkn1bz',
 'cdkn1ca',
 'cdkn1cb',
 'cdkn1cc',
 'cdkn1cd',
 'cdkn1ce',
 'cdkn1cf',
 'cdkn1cg',
 'cdkn1ch',
 'cdkn1ci',
 'cdkn1cj',
 'cdkn1ck',
 'cdkn1cl',
 'cdkn1cm',
 'cdkn1cn',
 'cdkn1co',
 'cdkn1cp',
 'cdkn1cq',
 'cdkn1cr',
 'cdkn1cs',
 'cdkn1ct',
 'cdkn1cu',
 'cdkn1cv',
 'cdkn1cw',
 'cdkn1cx',
 'cdkn1cy',
 'cdkn1cz',
 'cdkn1da',
 'cdkn1db',
 'cdkn1dc',
 'cdkn1dd',
 'cdkn1de',
 'cdkn1df',
 'cdkn1dg',
 'cdkn1dh',
 'cdkn1di',
 'cdkn1dj',
 'cdkn1dk',
 'cdkn1dl',
 'cdkn1dm',
 'cdkn1dn',
 'cdkn1do',
 'cdkn1dp',
 'cdkn1dq',
 'cdkn1dr',
 'cdkn1ds',
 'cdkn1dt',
 'cdkn1du',
 'cdkn1dv',
 'cdkn1dw',
 'cdkn1dx',
 'cdkn1dy',
 'cdkn1dz',
 'cdkn1ea',
 'cdkn1eb',
 'cdkn1ec',
 'cdkn1ed',
 'cdkn1ee',
 'cdkn1ef',
 'cdkn1eg',
 'cdkn1eh',
 'cdkn1ei',
 'cdkn1ej',
 'cdkn1ek',
 'cdkn1el',
 'cdkn1em',
 'cdkn1en',
 'cdkn1eo',
 'cdkn1ep',
 'cdkn1eq',
 'cdkn1er',
 'cdkn1es',
 'cdkn1et',
 'cdkn1eu',
 'cdkn1ev',
 'cdkn1ew',
 'cdkn1ex',
 'cdkn1ey',
 'cdkn1ez',
 'cdkn1fa',
 'cdkn1fb',
 'cdkn1fc',
 'cdkn1fd',
 'cdkn1fe',
 'cdkn1ff',
 'cdkn1fg',
 'cdkn1fh',
 'cdkn1fi',
 'cdkn1fj',
 'cdkn1fk',
 'cdkn1fl',
 'cdkn1fm',
 'cdkn1fn',
 'cdkn1fo',
 'cdkn1fp',
 'cdkn1fq',
 'cdkn1fr',
 'cdkn1fs',
 'cdkn1ft',
 'cdkn1fu',
 'cdkn1fv',
 'cdkn1fw',
 'cdkn1fx',
 'cdkn1fy',
 'cdkn1fz',
 'cdkn1ga',
 'cdkn1gb',
 'cdkn1gc',
 'cdkn1gd',
 'cdkn1ge',
 'cdkn1gf',
 'cdkn1gg',
 'cdkn1gh',
 'cdkn1gi',
 'cdkn1gj',
 'cdkn1gk',
 'cdkn1gl',
 'cdkn1gm',
 'cdkn1gn',
 'cdkn1go',
 'cdkn1gp',
 'cdkn1gq',
 'cdkn1gr',
 'cdkn1gs',
 'cdkn1gt',
 'cdkn1gu',
 'cdkn1gv',
 'cdkn1gw',
 'cdkn1gx',
 'cdkn1gy',
 'cdkn1gz',
 'cdkn1ha',
 'cdkn1hb',
 'cdkn1hc',
 'cdkn1hd',
 'cdkn1he',
 'cdkn1hf',
 'cdkn1hg',
 'cdkn1hh',
 'cdkn1hi',
 'cdkn1hj',
 'cdkn1hk',
 'cdkn1hl',
 'cdkn1hm',
 'cdkn1hn',
 'cdkn1ho',
 'cdkn1hp',
 'cdkn1hq',
 'cdkn1hr',
 'cdkn1hs',
 'cdkn1ht',
 'cdkn1hu',
 'cdkn1hv',
 'cdkn1hw',
 'cdkn1hx',
 'cdkn1hy',
 'cdkn1hz',
 'cdkn1ia',
 'cdkn1ib',
 'cdkn1ic',
 'cdkn1id',
 'cdkn1ie',
 'cdkn1if',
 'cdkn1ig',
 'cdkn1ih',
 'cdkn1ii',
 'cdkn1ij',
 'cdkn1ik',
 'cdkn1il',
 'cdkn1im',
 'cdkn1in',
 'cdkn1io',
 'cdkn1ip',
 'cdkn1iq',
 'cdkn1ir',
 'cdkn1is',
 'cdkn1it',
 'cdkn1iu',
 'cdkn1iv',
 'cdkn1iw',
 'cdkn1ix',
 'cdkn1iy',
 'cdkn1iz',
 'cdkn1ja',
 'cdkn1jb',
 'cdkn1jc',
 'cdkn1jd',
 'cdkn1je',
 'cdkn1jf',
 'cdkn1jg',
 'cdkn1jh',
 'cdkn1ji',
 'cdkn1jj',
 'cdkn1jk',
 'cdkn1jl',
 'cdkn1jm',
 'cdkn1jn',
 'cdkn1jo',
 'cdkn1jp',
 'cdkn1jq',
 'cdkn1jr',
 'cdkn1js',
 'cdkn1jt',
 'cdkn1ju',
 'cdkn1jv',
 'cdkn1jw',
 'cdkn1jx',
 'cdkn1jy',
 'cdkn1jz',
 'cdkn1ka',
 'cdkn1kb',
 'cdkn1kc',
 'cdkn1kd',
 'cdkn1ke',
 'cdkn1kf',
 'cdkn1kg',
 'cdkn1kh',
 'cdkn1ki',
 'cdkn1kj',
 'cdkn1kk',
 'cdkn1kl',
 'cdkn1km',
 'cdkn1kn',
 'cdkn1ko',
 'cdkn1kp',
 'cdkn1kq',
 'cdkn1kr',
 'cdkn1ks',
 'cdkn1kt',
 'cdkn1ku',
 'cdkn1kv',
 'cdkn1kw',
 'cdkn1kx',
 'cdkn1ky',
 'cdkn1kz',
 'cdkn1la',
 'cdkn1lb',
 'cdkn1lc',
 'cdkn1ld',
 'cdkn1le',
 'cdkn1lf',
 'cdkn1lg',
 'cdkn1lh',
 'cdkn1li',
 'cdkn1lj',
 'cdkn1lk',
 'cdkn1ll',
 'cdkn1lm',
 'cdkn1ln',
 'cdkn1lo',
 'cdkn1lp',
 'cdkn1lq',
 'cdkn1lr',
 'cdkn1ls',
 'cdkn1lt',
 'cdkn1lu',
 'cdkn1lv',
 'cdkn1lw',
 'cdkn1lx',
 'cdkn1ly',
 'cdkn1lz',
 'cdkn1ma',
 'cdkn1mb',
 'cdkn1mc',
 'cdkn1md',
 'cdkn1me',
 'cdkn1mf',
 'cdkn1mg',
 'cdkn1mh',
 'cdkn1mi',
 'cdkn1mj',
 'cdkn1mk',
 'cdkn1ml',
 'cdkn1mm',
 'cdkn1mn',
 'cdkn1mo',
 'cdkn1mp',
 'cdkn1mq',
 'cdkn1mr',
 'cdkn1ms',
 'cdkn1mt',
 'cdkn1mu',
 'cdkn1mv',
 'cdkn1mw',
 'cdkn1mx',
 'cdkn1my',
 'cdkn1mz',
 'cdkn1na',
 'cdkn1nb',
 'cdkn1nc',
 'cdkn1nd',
 'cdkn1ne',
 'cdkn1nf',
 'cdkn1ng',
 'cdkn1nh',
 'cdkn1ni',
 'cdkn1nj',
 'cdkn1nk',
 'cdkn1nl',
 'cdkn1nm',
 'cdkn1nn',
 'cdkn1no',
 'cdkn1np',
 'cdkn1nq',
 'cdkn1nr',
 'cdkn1ns',
 'cdkn1nt',
 'cdkn1nu',
 'cdkn1nv',
 'cdkn1nw',
 'cdkn1nx',
 'cdkn1ny',
 'cdkn1nz',
 'cdkn1oa',
 'cdkn1ob',
 'cdkn1oc',
 'cdkn1od',
 'cdkn1oe',
 'cdkn1of',
 'cdkn1og',
 'cdkn1oh',
 'cdkn1oi',
 'cdkn1oj',
 'cdkn1ok',
 'cdkn1ol',
 'cdkn1om',
 'cdkn1on',
 'cdkn1oo',
 'cdkn1op',
 'cdkn1oq',
 'cdkn1or',
 'cdkn1os',
 'cdkn1ot',
 'cdkn1ou',
 'cdkn1ov',
 'cdkn1ow',
 'cdkn1ox',
 'cdkn1oy',
 'cdkn1oz',
 'cdkn1pa',
 'cdkn1pb',
 'cdkn1pc',
 'cdkn1pd',
 'cdkn1pe',
 'cdkn1pf',
 'cdkn1pg',
 'cdkn1ph',
 'cdkn1pi',
 'cdkn1pj',
 'cdkn1pk',
 'cdkn1pl',
 'cdkn1pm',
 'cdkn1pn',
 'cdkn1po',
 'cdkn1pp',
 'cdkn1pq',
 'cdkn1pr',
 'cdkn1ps',
 'cdkn1pt',
 'cdkn1pu',
 'cdkn1pv',
 'cdkn1pw',
 'cdkn1px',
 'cdkn1py',
 'cdkn1pz',
 'cdkn1qa',
 'cdkn1qb',
 'cdkn1qc',
 'cdkn1qd',
 'cdkn1qe',
 'cdkn1qf',
 'cdkn1qg',
 'cdkn1qh',
 'cdkn1qi',
 'cdkn1qj',
 'cdkn1qk',
 'cdkn1ql',
 'cdkn1qm',
 'cdkn1qn',
 'cdkn1qo',
 'cdkn1qp',
 'cdkn1qq',
 'cdkn1qr',
 'cdkn1qs',
 'cdkn1qt',
 'cdkn1qu',
 'cdkn1qv',
 'cdkn1qw',
 'cdkn1qx',
 'cdkn1qy',
 'cdkn1qz',
 'cdkn1ra',
 'cdkn1rb',
 'cdkn1rc',
 'cdkn1rd',
 'cdkn1re',
 'cdkn1rf',
 'cdkn1rg',
 'cdkn1rh',
 'cdkn1ri',
 'cdkn1rj',
 'cdkn1rk',
 'cdkn1rl',
 'cdkn1rm',
 'cdkn1rn',
 'cdkn1ro',
 'cdkn1rp',
 'cdkn1rq',
 'cdkn1rr',
 'cdkn1rs',
 'cdkn1rt',
 'cdkn1ru',
 'cdkn1rv',
 'cdkn1rw',
 'cdkn1rx',
 'cdkn1ry',
 'cdkn1rz',
 'cdkn1sa',
 'cdkn1sb',
 'cdkn1sc',
 'cdkn1sd',
 'cdkn1se',
 'cdkn1sf',
 'cdkn1sg',
 'cdkn1sh',
 'cdkn1si',
 'cdkn1sj',
 'cdkn1sk',
 'cdkn1sl',
 'cdkn1sm',
 'cdkn1sn',
 'cdkn1so',
 'cdkn1sp',
 'cdkn1sq',
 'cdkn1sr',
 'cdkn1ss',
 'cdkn1st',
 'cdkn1su',
 'cdkn1sv',
 'cdkn1sw',
 'cdkn1sx',
 'cdkn1sy',
 'cdkn1sz',
 'cdkn1ta',
 'cdkn1tb',
 'cdkn1tc',
 'cdkn1td',
 'cdkn1te',
 'cdkn1tf',
 'cdkn1tg',
 'cdkn1th',
 'cdkn1ti',
 'cdkn1tj',
 'cdkn1tk',
 'cdkn1tl',
 'cdkn1tm',
 'cdkn1tn',
 'cdkn1to',
 'cdkn1tp',
 'cdkn1tq',
 'cdkn1tr',
 'cdkn1ts',
 'cdkn1tt',
 'cdkn1tu',
 'cdkn1tv',
 'cdkn1tw',
 'cdkn1tx',
 'cdkn1ty',
 'cdkn1tz',
 'cdkn1ua',
 'cdkn1ub',
 'cdkn1uc',
 'cdkn1ud',
 'cdkn1ue',
 'cdkn1uf',
 'cdkn1ug',
 'cdkn1uh',
 'cdkn1ui',
 'cdkn1uj',
 'cdkn1uk',
 'cdkn1ul',
 'cdkn1um',
 'cdkn1un',
 'cdkn1uo',
 'cdkn1up',
 'cdkn1uq',
 'cdkn1ur',
 'cdkn1us',
 'cdkn1ut',
 'cdkn1uu',
 'cdkn1uv',
 'cdkn1uw',
 'cdkn1ux',
 'cdkn1uy',
 'cdkn1uz',
 'cdkn1va',
 'cdkn1vb',
 'cdkn1vc',
 'cdkn1vd',
 'cdkn1ve',
 'cdkn1vf',
 'cdkn1vg',
 'cdkn1vh',
 'cdkn1vi',
 'cdkn1vj',
 'cdkn1vk',
 'cdkn1vl',
 'cdkn1vm',
 'cdkn1vn',
 'cdkn1vo',
 'cdkn1vp',
 'cdkn1vq',
 'cdkn1vr',
 'cdkn1vs',
 'cdkn1vt',
 'cdkn1vu',
 'cdkn1vv',
 'cdkn1vw',
 'cdkn1vx',
 'cdkn1vy',
 'cdkn1vz',
 'cdkn1wa',
 'cdkn1wb',
 'cdkn1wc',
 'cdkn1wd',
 'cdkn1we',
 'cdkn1wf',
 'cdkn1wg',
 'cdkn1wh',
 'cdkn1wi',
 'cdkn1wj',
 'cdkn1wk',
 'cdkn1wl',
 'cdkn1wm',
 'cdkn1wn',
 'cdkn1wo',
 'cdkn1wp',
 'cdkn1wq',
 'cdkn1wr',
 'cdkn1ws',
 'cdkn1wt',
 'cdkn1wu',
 'cdkn1wv',
 'cdkn1ww',
 'cdkn1wx',
 'cdkn1wy',
 'cdkn1wz',
 'cdkn1xa',
 'cdkn1xb',
 'cdkn1xc',
 'cdkn1xd',
 'cdkn1xe',
 'cdkn1xf',
 'cdkn1xg',
 'cdkn1xh',
 'cdkn1xi',
 'cdkn1xj',
 'cdkn1xk',
 'cdkn1xl',
 'cdkn1xm',
 'cdkn1xn',
 'cdkn1xo',
 'cdkn1xp',
 'cdkn1xq',
 'cdkn1xr',
 'cdkn1xs',
 'cdkn1xt',
 'cdkn1xu',
 'cdkn1xv',
 'cdkn1xw',
 'cdkn1xx',
 'cdkn1xy',
 'cdkn1xz',
 'cdkn1ya',
 'cdkn1yb',
 'cdkn1yc',
 'cdkn1yd',
 'cdkn1ye',
 'cdkn1yf',
 'cdkn1yg',
 'cdkn1yh',
 'cdkn1yi',
 'cdkn1yj',
 'cdkn1yk',
 'cdkn1yl',
 'cdkn1ym',
 'cdkn1yn',
 'cdkn1yo',
 'cdkn1yp',
 'cdkn1yq',
 'cdkn1yr',
 'cdkn1ys',
 'cdkn1yt',
 'cdkn1yu',
 'cdkn1yv',
 'cdkn1yw',
 'cdkn1yx',
 'cdkn1yy',
 'cdkn1yz',
 'cdkn1za',
 'cdkn1zb',
 'cdkn1zc',
 'cdkn1zd',
 'cdkn1ze',
 'cdkn1zf',
 'cdkn1zg',
 'cdkn1zh',
 'cdkn1zi',
 'cdkn1zj',
 'cdkn1zk',
 'cdkn1zl',
 'cdkn1zm',
 'cdkn1zn',
 'cdkn1zo',
 'cdkn1zp',
 'cdkn1zq',
 'cdkn1zr',
 'cdkn1zs',
 'cdkn1zt',
 'cdkn1zu',
 'cdkn1zv',
 'cdkn1zw',
 'cdkn1zx',
 'cdkn1zy',
 'cdkn1zz']
```

'cdkn1b',
'cdkn2a',
'cdkn2b',
'cdkn2c',
'chek2',
'cic',
'crebbp',
'ctcf',
'ctnnb1',
'ddr2',
'dicer1',
'dnmt3a',
'dnmt3b',
'egfr',
'eif1ax',
'elf3',
'ep300',
'epas1',
'epcam',
'erbb2',
'erbb3',
'erbb4',
'ercc2',
'ercc3',
'ercc4',
'erg',
'esr1',
'etv6',
'ewsr1',
'ezh2',
'fanca',
'fancc',
'fat1',
'fbxw7',
'fgf19',
'fgf4',
'fgfr1',
'fgfr2',
'fgfr3',
'fgfr4',
'flt3',
'foxa1',
'foxo1',
'foxp1',
'fubp1',
'gata3',
'gnaq',
'gnas',
'h3f3a',
'hist1h1c',
'hla',
'hnf1a',
'hras',
'idh1',
'idh2',
'igf1r',
'ikbke',
'ikzf1',
'jak1',
'jak2',
'jun',
'kdm5a',
'kdm5c',
'kdm6a',
'kdr',
'keap1',
'kit',
'kmt2a',
'kmt2b',
'kmt2c',
'kmt2d',
'knstrn',
'kras',
'lats1',
'map2k1',
'map2k2',
'map2k4',

'map3k1',
'mdm2',
'mdm4',
'med12',
'mef2b',
'men1',
'met',
'mga',
'mlh1',
'mpl',
'msh2',
'msh6',
'mtor',
'myc',
'mycn',
'myd88',
'myod1',
'ncor1',
'nfl1',
'nf2',
'nfe2l2',
'nfkb1a',
'nkx2',
'notch1',
'notch2',
'npm1',
'nras',
'nsd1',
'ntrk1',
'ntrk2',
'ntrk3',
'nup93',
'pak1',
'pbrm1',
'pdgfra',
'pdgfrb',
'pik3ca',
'pik3cb',
'pik3cd',
'pik3r1',
'pik3r2',
'pik3r3',
'pim1',
'pms2',
'pole',
'ppm1d',
'ppp2r1a',
'ppp6c',
'prdm1',
'ptch1',
'pten',
'ptpn11',
'ptprd',
'ptprt',
'rab35',
'rac1',
'rad21',
'rad50',
'rad51c',
'rad51d',
'rad54l',
'raf1',
'rasa1',
'rb1',
'rbm10',
'ret',
'rheb',
'rhoa',
'rit1',
'rnf43',
'ros1',
'rras2',
'runx1',
'rxra',
'sdhb',
'sdhc',
'setd2',

```
'sf3b1',
'shoc2',
'smad2',
'smad3',
'smad4',
'smarca4',
'smarcb1',
'smo',
'sos1',
'sox9',
'spop',
'src',
'stag2',
'stat3',
'stk11',
'tcf3',
'tert',
'tet1',
'tet2',
'tgfbr1',
'tgfbr2',
'tmprss2',
'tp53',
'tp53bp1',
'tsc1',
'tsc2',
'u2af1',
'vegfa',
'vhl',
'whsc1',
'xpo1',
'yap1']
```

In [26]:

```
print("train_gene_feature_onehotCoding is converted feature using one-hot encoding method. The shape of gene feature:", train_gene_feature_onehotCoding.shape)
```

train_gene_feature_onehotCoding is converted feature using one-hot encoding method. The shape of gene feature: (2124, 230)

Q4. How good is this gene feature in predicting y_i ?

There are many ways to estimate how good a feature is, in predicting y_i . One of the good methods is to build a proper ML model using just this feature. In this case, we will build a logistic regression model using only Gene feature (one hot encoded) to predict y_i .

In [27]:

```
alpha = [10 ** x for x in range(-5, 1)] # hyperparam for SGD classifier.

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link:
#-----

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_gene_feature_onehotCoding, y_train)
```

```

sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_gene_feature_onehotCoding, y_train)
predict_y = sig_clf.predict_proba(cv_gene_feature_onehotCoding)
cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))


fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_gene_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_gene_feature_onehotCoding, y_train)

predict_y = sig_clf.predict_proba(train_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

For values of alpha = 1e-05 The log loss is: 1.3931649870541356
 For values of alpha = 0.0001 The log loss is: 1.2123813661174583
 For values of alpha = 0.001 The log loss is: 1.223407911806416
 For values of alpha = 0.01 The log loss is: 1.3504160696302134
 For values of alpha = 0.1 The log loss is: 1.4584693513489653
 For values of alpha = 1 The log loss is: 1.4904158719525435



For values of best alpha = 0.0001 The train log loss is: 1.0469768840690201
 For values of best alpha = 0.0001 The cross validation log loss is: 1.2123813661174583
 For values of best alpha = 0.0001 The test log loss is: 1.200092586520076

Q5. Is the Gene feature stable across all the data sets (Test, Train, Cross validation)?

Ans. Yes, it is. Otherwise, the CV and Test errors would be significantly more than train error.

In [28]:

```

print("Q6. How many data points in Test and CV datasets are covered by the ", unique_genes.shape[0], " genes in train dataset?")

test_coverage=test_df[test_df['Gene'].isin(list(set(train_df['Gene'])))].shape[0]
cv_coverage=cv_df[cv_df['Gene'].isin(list(set(train_df['Gene'])))].shape[0]

```



```
print('Ans\n1. In test data',test_coverage, 'out of',test_df.shape[0], ":", (test_coverage/test_df.shape[0])*100)
print('2. In cross validation data',cv_coverage, 'out of ',cv_df.shape[0]," :", (cv_coverage/cv_df.shape[0])*100)
```

Q6. How many data points in Test and CV datasets are covered by the 230 genes in train dataset?

Ans

1. In test data 640 out of 665 : 96.2406015037594
2. In cross validation data 517 out of 532 : 97.18045112781954

3.2.2 Univariate Analysis on Variation Feature

Q7. Variation, What type of feature is it ?

Ans. Variation is a categorical variable

Q8. How many categories are there?

In [29]:

```
unique_variations = train_df['Variation'].value_counts()
print('Number of Unique Variations :', unique_variations.shape[0])
# the top 10 variations that occurred most
print(unique_variations.head(10))
```

```
Number of Unique Variations : 1930
Truncating_Mutations    59
Amplification            47
Deletion                 47
Fusions                  21
Overexpression           5
E17K                     3
G12V                     3
S308A                    2
Q61R                     2
P34R                     2
Name: Variation, dtype: int64
```


In [30]:


```
print("Ans: There are", unique_variations.shape[0] ,"different categories of variations in the train data, and they are distributed as follows",)
```

Ans: There are 1930 different categories of variations in the train data, and they are distributed as follows

In [31]:

```
s = sum(unique_variations.values);
h = unique_variations.values/s;
plt.plot(h, label="Histogram of Variations")
plt.xlabel('Index of a Variation')
plt.ylabel('Number of Occurances')
plt.legend()
plt.grid()
plt.show()
```






In [32]:

```
c = np.cumsum(h)
print(c)
plt.plot(c, label='Cumulative distribution of Variations')
plt.grid()
plt.legend()
plt.show()
```

```
[0.02777778 0.04990584 0.0720339 ... 0.99905838 0.99952919 1. ]
```



Q9. How to featurize this Variation feature ?

Ans. There are two ways we can featurize this variable check out this video:

<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>

1. One hot Encoding
2. Response coding

We will be using both these methods to featurize the Variation Feature

In [33]:

```
# alpha is used for laplace smoothing
alpha = 1
# train gene feature
train_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", train_df))
# test gene feature
test_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", test_df))
# cross validation gene feature
cv_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", cv_df))
```

In [34]:

```
print("train_variation_feature_responseCoding is a converted feature using the response coding method. The shape of Variation feature:", train_variation_feature_responseCoding.shape)
```

train_variation_feature_responseCoding is a converted feature using the response coding method. The shape of Variation feature: (2124, 9)

In [35]:

```
# one-hot encoding of variation feature.
variation_vectorizer = CountVectorizer()
train_variation_feature_onehotCoding = variation_vectorizer.fit_transform(train_df['Variation'])
test_variation_feature_onehotCoding = variation_vectorizer.transform(test_df['Variation'])
cv_variation_feature_onehotCoding = variation_vectorizer.transform(cv_df['Variation'])
```

In [36]:

```
print("train_variation_feature_onehotEncoded is converted feature using the onne-hot encoding meth  
od. The shape of Variation feature:", train_variation_feature_onehotCoding.shape)
```

train_variation_feature_onehotEncoded is converted feature using the onne-hot encoding method. The shape of Variation feature: (2124, 1962)

Q10. How good is this Variation feature in predicting y_i?

Let's build a model just like the earlier!


In [37]:

```
alpha = [10 ** x for x in range(-5, 1)]

# read more about SGDClassifier() at http://scikit-  
learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html  
# -----  
# default parameters  
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_i  
ter=None, tol=None,  
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0  
=0.0, power_t=0.5,  
# class_weight=None, warm_start=False, average=False, n_iter=None)  
  
# some of methods  
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.  
# predict(X) Predict class labels for samples in X.  
  
#-----  
# video link:  
#-----  
  
cv_log_error_array=[]  
for i in alpha:  
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)  
    clf.fit(train_variation_feature_onehotCoding, y_train)  
  
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")  
    sig_clf.fit(train_variation_feature_onehotCoding, y_train)  
    predict_y = sig_clf.predict_proba(cv_variation_feature_onehotCoding)  
  
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))  
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.clas  
ses_, eps=1e-15))  
  
fig, ax = plt.subplots()  
ax.plot(alpha, cv_log_error_array, c='g')  
for i, txt in enumerate(np.round(cv_log_error_array, 3)):  
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))  
plt.grid()  
plt.title("Cross Validation Error for each alpha")  
plt.xlabel("Alpha i's")  
plt.ylabel("Error measure")  
plt.show()  
  
best_alpha = np.argmin(cv_log_error_array)  
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)  
clf.fit(train_variation_feature_onehotCoding, y_train)  
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")  
sig_clf.fit(train_variation_feature_onehotCoding, y_train)  
  
predict_y = sig_clf.predict_proba(train_variation_feature_onehotCoding)  
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train,  
predict_y, labels=clf.classes_, eps=1e-15))  
predict_y = sig_clf.predict_proba(cv_variation_feature_onehotCoding)  
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_lo  
ss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))  
predict_y = sig_clf.predict_proba(test_variation_feature_onehotCoding)
```

```
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
```

For values of alpha = 1e-05 The log loss is: 1.7534874632262567
 For values of alpha = 0.0001 The log loss is: 1.7425754054760152
 For values of alpha = 0.001 The log loss is: 1.7407628451858137
 For values of alpha = 0.01 The log loss is: 1.748631223517745
 For values of alpha = 0.1 The log loss is: 1.7573081676774112
 For values of alpha = 1 The log loss is: 1.7561090309386373



For values of best alpha = 0.001 The train log loss is: 1.0669439311456297
 For values of best alpha = 0.001 The cross validation log loss is: 1.7407628451858137
 For values of best alpha = 0.001 The test log loss is: 1.6997242940202788

Q11. Is the Variation feature stable across all the data sets (Test, Train, Cross validation)?

Ans. Not sure! But lets be very sure using the below analysis.

In [38]:

```
print("Q12. How many data points are covered by total ", unique_variations.shape[0], " genes in test and cross validation data sets?")
test_coverage=test_df[test_df['Variation'].isin(list(set(train_df['Variation'])))].shape[0]
cv_coverage=cv_df[cv_df['Variation'].isin(list(set(train_df['Variation'])))].shape[0]
print('Ans\1. In test data',test_coverage, 'out of',test_df.shape[0], ":", (test_coverage/test_df.shape[0])*100)
print('2. In cross validation data',cv_coverage, 'out of ',cv_df.shape[0],":", (cv_coverage/cv_df.shape[0])*100)
```

Q12. How many data points are covered by total 1930 genes in test and cross validation data sets?

Ans

1. In test data 77 out of 665 : 11.578947368421053
2. In cross validation data 48 out of 532 : 9.022556390977442

3.2.3 Univariate Analysis on Text Feature

1. How many unique words are present in train data?
2. How are word frequencies distributed?
3. How to featurize text field?
4. Is the text feature useful in predicting y_i?
5. Is the text feature stable across train, test and CV datasets?

In [39]:

```
# cls_text is a data frame
# for every row in data fram consider the 'TEXT'
# split the words by space
# make a dict with those words
# increment its count whenever we see that word
```

```
def extract_dictionary_paddle(cls_text):
    dictionary = defaultdict(int)
    for index, row in cls_text.iterrows():
        for word in row['TEXT'].split():
            dictionary[word] +=1
    return dictionary
```

In [40]:

```
import math
#https://stackoverflow.com/a/1602964
def get_text_responsecoding(df):
    text_feature_responseCoding = np.zeros((df.shape[0],9))
    for i in range(0,9):
        row_index = 0
        for index, row in df.iterrows():
            sum_prob = 0
            for word in row['TEXT'].split():
                sum_prob += math.log(((dict_list[i].get(word,0)+10)/(total_dict.get(word,0)+90)))
            text_feature_responseCoding[row_index][i] = math.exp(sum_prob/len(row['TEXT'].split()))
            row_index += 1
    return text_feature_responseCoding
```

Considering top 1000 features as per idf values

In [41]:

```
# building a tfidfvectorizer with all the words that occurred minimum 3 times in train data
text_vectorizer = TfidfVectorizer(min_df=3)
train_text_feature_tfidf = text_vectorizer.fit_transform(train_df['TEXT'])

# getting all the feature names (words)
tfidf_train_text_onehotencoding= text_vectorizer.get_feature_names()

# creating dictionary
dictionary = dict(zip(text_vectorizer.get_feature_names(), list(text_vectorizer.idf_)))

from collections import OrderedDict
sorted_by_value = OrderedDict(sorted(dictionary.items(),reverse= True, key=lambda x: x[1]))

# getting top 1k features using idf_ values
top_tfidf_features = list(sorted_by_value.keys())[:1000]
```

In [42]:

```
# building a tfidfvectorizer with top 1k words that occurred minimum 3 times in train data
text_vectorizer = TfidfVectorizer(min_df=3,vocabulary=top_tfidf_features)
text_vectorizer.fit(train_df['TEXT'])
train_text_feature_onehotCoding = text_vectorizer.transform(train_df['TEXT'])

# getting all the feature names (words)
tfidf_train_text_onehotencoding= text_vectorizer.get_feature_names()
print("len of feature names is {}".format(len(tfidf_train_text_onehotencoding)))

# train_text_feature_onehotCoding.sum(axis=0).A1 will sum every row and returns (1*number of features) vector
train_text_fea_counts = train_text_feature_onehotCoding.sum(axis=0).A1

# zip(list(text_features),text_fea_counts) will zip a word with its number of times it occurred
text_fea_dict = dict(zip(list(tfidf_train_text_onehotencoding),train_text_fea_counts))
```

len of feature names is 1000

In [43]:

```
dict_list = []
# dict_list =[] contains 9 dictionaries each corresponds to a class
for i in range(1,10):
    cls_text = train_df[train_df['Class']==i]
```

```

cls_conf = train_df[train_df['class'] == cls]
# build a word dict based on the words in that class
dict_list.append(extract_dictionary_paddle(cls_text))
# append it to dict_list

# dict_list[i] is build on i'th class text data
# total_dict is build on whole training text data
total_dict = extract_dictionary_paddle(train_df)

confuse_array = []
for i in tfidf_train_text_onehotencoding:
    ratios = []
    max_val = -1
    for j in range(0,9):
        ratios.append((dict_list[j][i]+10)/(total_dict[i]+90))
    confuse_array.append(ratios)
confuse_array = np.array(confuse_array)

```

In [44]:

```

#response coding of text features
train_text_feature_responseCoding = get_text_responsecoding(train_df)
test_text_feature_responseCoding = get_text_responsecoding(test_df)
cv_text_feature_responseCoding = get_text_responsecoding(cv_df)

```

In [45]:

```

# https://stackoverflow.com/a/16202486
# we convert each row values such that they sum to 1
train_text_feature_responseCoding =
(train_text_feature_responseCoding.T/train_text_feature_responseCoding.sum(axis=1)).T
test_text_feature_responseCoding =
(test_text_feature_responseCoding.T/test_text_feature_responseCoding.sum(axis=1)).T
cv_text_feature_responseCoding = (cv_text_feature_responseCoding.T/cv_text_feature_responseCoding.
sum(axis=1)).T

```

In [46]:

```

# don't forget to normalize every feature
train_text_feature_onehotCoding = normalize(train_text_feature_onehotCoding, axis=0)

# we use the same vectorizer that was trained on train data
test_text_feature_onehotCoding = text_vectorizer.transform(test_df['TEXT'])
# don't forget to normalize every feature
test_text_feature_onehotCoding = normalize(test_text_feature_onehotCoding, axis=0)

# we use the same vectorizer that was trained on train data
cv_text_feature_onehotCoding = text_vectorizer.transform(cv_df['TEXT'])
# don't forget to normalize every feature
cv_text_feature_onehotCoding = normalize(cv_text_feature_onehotCoding, axis=0)

```

In [47]:

```

#https://stackoverflow.com/a/2258273/4084039
sorted_text_fea_dict = dict(sorted(text_fea_dict.items(), key=lambda x: x[1] , reverse=True))
sorted_text_occur = np.array(list(sorted_text_fea_dict.values()))

```

In [48]:

```

# Train a Logistic regression+Calibration model using text features which are on-hot encoded
alpha = [10 ** x for x in range(-5, 1)]

# read more about SGDClassifier() at http://scikit-
learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_i
ter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0
=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

```

```

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link:
#-----

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_text_feature_onehotCoding, y_train)

    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_text_feature_onehotCoding, y_train)
    predict_y = sig_clf.predict_proba(cv_text_feature_onehotCoding)
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))


fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_text_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_text_feature_onehotCoding, y_train)

predict_y = sig_clf.predict_proba(train_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

For values of alpha = 1e-05 The log loss is: 1.7534945992767754
 For values of alpha = 0.0001 The log loss is: 1.737408750096627
 For values of alpha = 0.001 The log loss is: 1.745960801103881
 For values of alpha = 0.01 The log loss is: 1.7489726484969215
 For values of alpha = 0.1 The log loss is: 1.7541927730548152
 For values of alpha = 1 The log loss is: 1.7564187975785703



For values of best alpha = 0.0001 The train log loss is: 1.593989512639425
 For values of best alpha = 0.0001 The cross validation log loss is: 1.737408750096627
 For values of best alpha = 0.0001 The test log loss is: 1.7100000000000001

for values of best alpha = 0.0001 the test log loss is: 1.122239822740861

Q. Is the Text feature stable across all the data sets (Test, Train, Cross validation)?

Ans. Yes, it seems like!

In [49]:

```
def get_intersec_text(df):
    df_text_vec = CountVectorizer(min_df=3)
    df_text_fea = df_text_vec.fit_transform(df['TEXT'])
    df_text_features = df_text_vec.get_feature_names()

    df_text_fea_counts = df_text_fea.sum(axis=0).A1
    df_text_fea_dict = dict(zip(list(df_text_features), df_text_fea_counts))
    len1 = len(set(df_text_features))
    len2 = len(set(tfidf_train_text_onehotencoding) & set(df_text_features))
    return len1, len2
```

In [50]:

```
len1, len2 = get_intersec_text(test_df)
print(np.round((len2/len1)*100, 3), "% of word of test data appeared in train data")
len1, len2 = get_intersec_text(cv_df)
print(np.round((len2/len1)*100, 3), "% of word of Cross Validation appeared in train data")
```

0.375 % of word of test data appeared in train data
0.052 % of word of Cross Validation appeared in train data

4. Machine Learning Models

In [51]:

```
#Data preparation for ML models.

#Misc. functionns for ML models

def predict_and_plot_confusion_matrix(train_x, train_y, test_x, test_y, clf):
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)
    pred_y = sig_clf.predict(test_x)

    # for calculating log_loss we will provide the array of probabilities belongs to each class
    print("Log loss :", log_loss(test_y, sig_clf.predict_proba(test_x)))
    # calculating the number of data points that are misclassified
    mis_classified = np.count_nonzero((pred_y - test_y))/test_y.shape[0]
    print("Number of mis-classified points :", mis_classified)
    plot_confusion_matrix(test_y, pred_y)
    return mis_classified
```

In [52]:

```
def report_log_loss(train_x, train_y, test_x, test_y, clf):
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)
    sig_clf_probs = sig_clf.predict_proba(test_x)
    return log_loss(test_y, sig_clf_probs, eps=1e-15)
```

In [53]:

```
# this function will be used just for naive bayes
# for the given indices, we will print the name of the features
# and we will check whether the feature present in the test point text or not
def get_impfeature_names(indices, text, gene, var, no_features):
    gene_count_vec = CountVectorizer()
    var_count_vec = CountVectorizer()
```



```

text_count_vec = CountVectorizer(min_df=3)

gene_vec = gene_count_vec.fit(train_df['Gene'])
var_vec = var_count_vec.fit(train_df['Variation'])
text_vec = text_count_vec.fit(train_df['TEXT'])

fea1_len = len(gene_vec.get_feature_names())
fea2_len = len(var_count_vec.get_feature_names())

word_present = 0
for i,v in enumerate(indices):
    if (v < fea1_len):
        word = gene_vec.get_feature_names()[v]
        yes_no = True if word == gene else False
        if yes_no:
            word_present += 1
            print(i, "Gene feature [{}] present in test data point [{}].".format(word,yes_no))
    elif (v < fea1_len+fea2_len):
        word = var_vec.get_feature_names()[v-(fea1_len)]
        yes_no = True if word == var else False
        if yes_no:
            word_present += 1
            print(i, "variation feature [{}] present in test data point [{}].".format(word,yes_r
o))
    else:
        word = text_vec.get_feature_names()[v-(fea1_len+fea2_len)]
        yes_no = True if word in text.split() else False
        if yes_no:
            word_present += 1
            print(i, "Text feature [{}] present in test data point [{}].".format(word,yes_no))

print("Out of the top ",no_features," features ", word_present, "are present in query point")

```

Stacking the three types of features

In [54]:

```

# merging gene, variance and text features

# building train, test and cross validation data sets
# a = [[1, 2],
#       [3, 4]]
# b = [[4, 5],
#       [6, 7]]
# hstack(a, b) = [[1, 2, 4, 5],
#                 [ 3, 4, 6, 7]]

train_gene_var_onehotCoding =
hstack((train_gene_feature_onehotCoding,train_variation_feature_onehotCoding))
test_gene_var_onehotCoding =
hstack((test_gene_feature_onehotCoding,test_variation_feature_onehotCoding))
cv_gene_var_onehotCoding = hstack((cv_gene_feature_onehotCoding,cv_variation_feature_onehotCoding)
)

train_x_onehotCoding = hstack((train_gene_var_onehotCoding, train_text_feature_onehotCoding)).tocsr()
train_y = np.array(list(train_df['Class']))

test_x_onehotCoding = hstack((test_gene_var_onehotCoding, test_text_feature_onehotCoding)).tocsr()
test_y = np.array(list(test_df['Class']))

cv_x_onehotCoding = hstack((cv_gene_var_onehotCoding, cv_text_feature_onehotCoding)).tocsr()
cv_y = np.array(list(cv_df['Class']))

train_gene_var_responseCoding =
np.hstack((train_gene_feature_responseCoding,train_variation_feature_responseCoding))
test_gene_var_responseCoding =
np.hstack((test_gene_feature_responseCoding,test_variation_feature_responseCoding))
cv_gene_var_responseCoding =
np.hstack((cv_gene_feature_responseCoding,cv_variation_feature_responseCoding))

train_x_responseCoding = np.hstack((train_gene_var_responseCoding,

```

```
train_text_feature_responseCoding))
test_x_responseCoding = np.hstack((test_gene_var_responseCoding, test_text_feature_responseCoding)
)
cv_x_responseCoding = np.hstack((cv_gene_var_responseCoding, cv_text_feature_responseCoding))
```

In [55]:

```
print("One hot encoding features :")
print("(number of data points * number of features) in train data = ", train_x_onehotCoding.shape)
print("(number of data points * number of features) in test data = ", test_x_onehotCoding.shape)
print("(number of data points * number of features) in cross validation data =", cv_x_onehotCoding
.shape)
```

```
One hot encoding features :
(number of data points * number of features) in train data = (2124, 3192)
(number of data points * number of features) in test data = (665, 3192)
(number of data points * number of features) in cross validation data = (532, 3192)
```

In [56]:

```
print(" Response encoding features :")
print("(number of data points * number of features) in train data = ", train_x_responseCoding.shap
e)
print("(number of data points * number of features) in test data = ", test_x_responseCoding.shape)
print("(number of data points * number of features) in cross validation data =",
cv_x_responseCoding.shape)
```

```
Response encoding features :
(number of data points * number of features) in train data = (2124, 27)
(number of data points * number of features) in test data = (665, 27)
(number of data points * number of features) in cross validation data = (532, 27)
```

Task 2 - Top 1000 of tf-idf values

4.1. Base Line Model

4.1.1. Naive Bayes

4.1.1.1. Hyper parameter tuning

In [57]:

```
# find more about Multinomial Naive base function here http://scikit-learn.org/stable/modules/generated/sklearn.naive\_bayes.MultinomialNB.html
# -----
# default paramters
# sklearn.naive_bayes.MultinomialNB(alpha=1.0, fit_prior=True, class_prior=None)

# some of methods of MultinomialNB()
# fit(X, y[, sample_weight]) Fit Naive Bayes classifier according to X, y
# predict(X) Perform classification on an array of test vectors X.
# predict_log_proba(X) Return log-probability estimates for the test vector X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/naive-bayes-algorithm-1/
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
```

```

# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/naive-bayes-algorithm-1/
# -----

alpha = [0.00001, 0.0001, 0.001, 0.1, 1, 10, 100, 1000]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = MultinomialNB(alpha=i)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability estimates
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(np.log10(alpha), cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (np.log10(alpha[i]), cv_log_error_array[i]))
plt.grid()
plt.xticks(np.log10(alpha))
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = MultinomialNB(alpha=alpha[best_alpha])
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

# summarizing data
nb_best_alpha = alpha[best_alpha]
nb_encoding = "One hot"

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
nb_train_log_loss = log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train,
predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
nb_cv_log_loss = log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_lo
ss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(test_x_onehotCoding)
nb_test_log_loss = log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, p
redict_y, labels=clf.classes_, eps=1e-15))


```

```

for alpha = 1e-05
Log Loss : 1.3136554465889085
for alpha = 0.0001
Log Loss : 1.3080456375047507
for alpha = 0.001
Log Loss : 1.3003343772618179
for alpha = 0.1
Log Loss : 1.2430403343832006
for alpha = 1
Log Loss : 1.2087502830555656
for alpha = 10
Log Loss : 1.2691679498304693
for alpha = 100
Log Loss : 1.3253687370120735
for alpha = 1000

```

Log Loss : 1.3345391215046745



For values of best alpha = 1 The train log loss is: 0.7777593935880002

For values of best alpha = 1 The cross validation log loss is: 1.2087502830555656

For values of best alpha = 1 The test log loss is: 1.163468237924221

4.1.1.2. Testing the model with best hyper paramters

In [58]:

```
# find more about Multinomial Naive base function here http://scikit-learn.org/stable/modules/generated/sklearn.naive\_bayes.MultinomialNB.html
# -----
# default paramters
# sklearn.naive_bayes.MultinomialNB(alpha=1.0, fit_prior=True, class_prior=None)

# some of methods of MultinomialNB()
# fit(X, y[, sample_weight]) Fit Naive Bayes classifier according to X, y
# predict(X) Perform classification on an array of test vectors X.
# predict_log_proba(X) Return log-probability estimates for the test vector X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/naive-bayes-algorithm-1/
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
# -----


clf = MultinomialNB(alpha=alpha[best_alpha])
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)
sig_clf.probs = sig_clf.predict_proba(cv_x_onehotCoding)
# to avoid rounding error while multiplying probabilites we use log-probability estimates
print("Log Loss :", log_loss(cv_y, sig_clf.probs))

nb_misclassified = np.count_nonzero((sig_clf.predict(cv_x_onehotCoding) - cv_y)) / cv_y.shape[0]
print("Number of missclassified point :", nb_misclassified)
plot_confusion_matrix(cv_y, sig_clf.predict(cv_x_onehotCoding.toarray()))
```

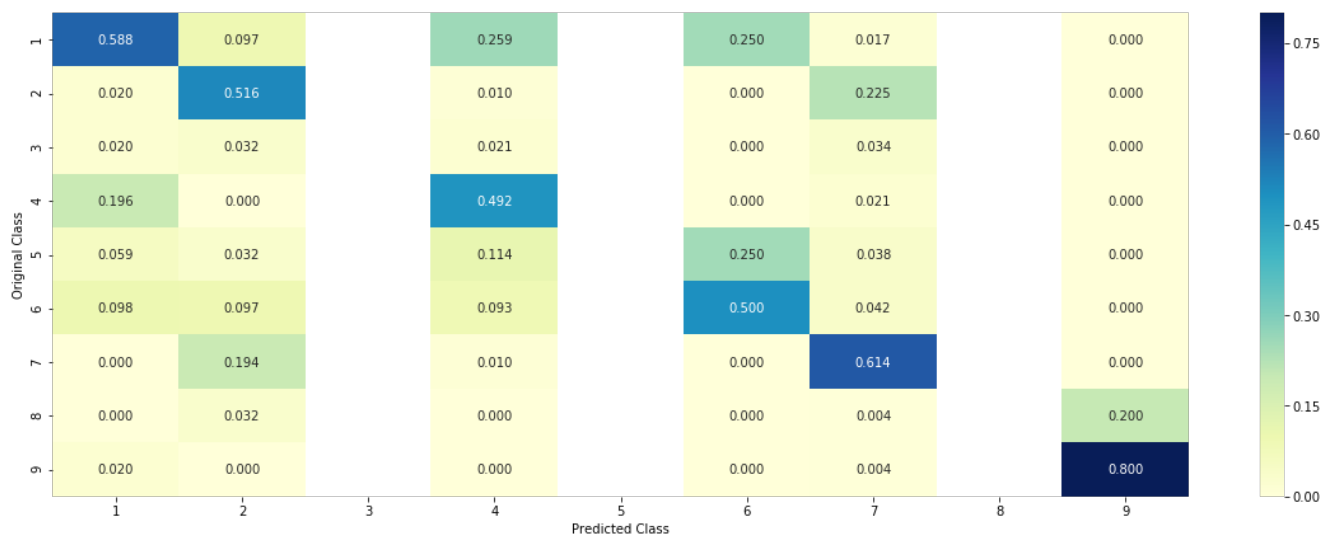
Log Loss : 1.2087502830555656

Number of missclassified point : 0.4398496240601504


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



4.1.1.3. Feature Importance, Correctly classified point

In [59]:

```
test_point_index = 1
```

```

no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index],test_df['Gene'].iloc[test_point_index],test_df['Variation']
.iloc[test_point_index], no_feature)

```

```

Predicted Class : 1
Predicted Class Probabilities: [[0.3049 0.1195 0.0473 0.1127 0.0827 0.0733 0.2392 0.0094 0.0109]]
Actual Class : 2
-----
Out of the top 100 features 0 are present in query point

```

4.1.1.4. Feature Importance, Incorrectly classified point

In [60]:

```

test_point_index = 100
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index],test_df['Gene'].iloc[test_point_index],test_df['Variation']
.iloc[test_point_index], no_feature)

```

```

Predicted Class : 4
Predicted Class Probabilities: [[0.3285 0.0606 0.0151 0.4274 0.0579 0.0415 0.0604 0.003 0.0055]]
Actual Class : 4
-----
98 Text feature [112] present in test data point [True]
Out of the top 100 features 1 are present in query point

```

4.2. K Nearest Neighbour Classification

4.2.1. Hyper parameter tuning

In [61]:

```

# find more about KNeighborsClassifier() here http://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html
# -----
# default parameter
# KNeighborsClassifier(n_neighbors=5, weights='uniform', algorithm='auto', leaf_size=30, p=2,
# metric='minkowski', metric_params=None, n_jobs=1, **kwargs)

# methods of
# fit(X, y) : Fit the model using X as training data and y as target values
# predict(X):Predict the class labels for the provided data
# predict_proba(X):Return probability estimates for the test data X.
#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/k-nearest-neighbors-geometric-intuition-with-a-toy-example-1/
#-----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)

```

```

#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [5, 11, 15, 21, 31, 41, 51, 99]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = KNeighborsClassifier(n_neighbors=i)
    clf.fit(train_x_responseCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_responseCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_responseCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability estimates
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

# summarizing data
knn_best_alpha = alpha[best_alpha]
knn_encoding = "Response"

predict_y = sig_clf.predict_proba(train_x_responseCoding)
knn_train_log_loss = log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train,
predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(cv_x_responseCoding)
knn_cv_log_loss = log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_lo
ss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(test_x_responseCoding)
knn_test_log_loss = log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, p
redict_y, labels=clf.classes_, eps=1e-15))


```

```

for alpha = 5
Log Loss : 1.045358515990766
for alpha = 11
Log Loss : 1.0500329323572986
for alpha = 15
Log Loss : 1.0538468137530441
for alpha = 21
Log Loss : 1.0615486777402305
for alpha = 31
Log Loss : 1.0643147940052855
for alpha = 41
Log Loss : 1.0735724357769685
for alpha = 51
Log Loss : 1.0760320934214398
for alpha = 99

```

Log Loss : 1.1063307306885015



For values of best alpha = 5 The train log loss is: 0.4868513233524239

For values of best alpha = 5 The cross validation log loss is: 1.045358515990766

For values of best alpha = 5 The test log loss is: 1.0314985157832384

4.2.2. Testing the model with best hyper paramters

In [62]:


```
# find more about KNeighborsClassifier() here http://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html
# -----
# default parameter
# KNeighborsClassifier(n_neighbors=5, weights='uniform', algorithm='auto', leaf_size=30, p=2,
# metric='minkowski', metric_params=None, n_jobs=1, **kwargs)

# methods of
# fit(X, y) : Fit the model using X as training data and y as target values
# predict(X):Predict the class labels for the provided data
# predict_proba(X):Return probability estimates for the test data X.
#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/k-nearest-neighbors-geometric-intuition-with-a-toy-example-1/
#-----
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
knn_misclassified = predict_and_plot_confusion_matrix(train_x_responseCoding, train_y,
cv_x_responseCoding, cv_y, clf)
```

Log loss : 1.045358515990766

Number of mis-classified points : 0.35150375939849626


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



4.2.3. Sample Query point -1

In [63]:

```

clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

test_point_index = 1
predicted_cls = sig_clf.predict(test_x_responseCoding[0].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Actual Class :", test_y[test_point_index])
neighbors = clf.kneighbors(test_x_responseCoding[test_point_index].reshape(1, -1), alpha[best_alpha])
print("The ", alpha[best_alpha], " nearest neighbours of the test points belongs to classes", train_y[neighbors[1][0]])
print("Fequency of nearest points :", Counter(train_y[neighbors[1][0]]))

```

Predicted Class : 4

Actual Class : 2

The 5 nearest neighbours of the test points belongs to classes [1 4 1 1 1]

Fequency of nearest points : Counter({1: 4, 4: 1})

4.2.4. Sample Query Point-2

In [64]:

```
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

test_point_index = 100

predicted_cls = sig_clf.predict(test_x_responseCoding[test_point_index].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Actual Class :", test_y[test_point_index])
neighbors = clf.kneighbors(test_x_responseCoding[test_point_index].reshape(1, -1), alpha[best_alpha])
print("the k value for knn is",alpha[best_alpha],"and the nearest neighbours of the test points belongs to classes",train_y[neighbors[1][0]])
print("Frequency of nearest points :",Counter(train_y[neighbors[1][0]]))
```

```
Predicted Class : 1
Actual Class : 4
the k value for knn is 5 and the nearest neighbours of the test points belongs to classes [1 1 1 1 4]
Frequency of nearest points : Counter({1: 4, 4: 1})
```

4.3. Logistic Regression

4.3.1. With Class balancing

4.3.1.1. Hyper parameter tuning

In [65]:

```
# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/geometric-intuition-1/
#-----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default parameters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10 ** x for x in range(-6, 3)]
cv_log_error_array = []
```

```

for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log', random_state=42)

    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability estimates
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

# summarizing data
lr_bal_best_alpha = alpha[best_alpha]
lr_bal_encoding = "One hot"

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
lr_bal_train_log_loss = log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
lr_bal_cv_log_loss = log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))


predict_y = sig_clf.predict_proba(test_x_onehotCoding)
lr_bal_test_log_loss = log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))


```

```

for alpha = 1e-06
Log Loss : 1.4813448288451678
for alpha = 1e-05
Log Loss : 1.3237210425671548
for alpha = 0.0001
Log Loss : 1.1526650264043543
for alpha = 0.001
Log Loss : 1.1715474156094234
for alpha = 0.01
Log Loss : 1.302319281232571
for alpha = 0.1
Log Loss : 1.4090101088261382
for alpha = 1
Log Loss : 1.4332908834114608
for alpha = 10
Log Loss : 1.437760865776187
for alpha = 100
Log Loss : 1.438349273548111

```





For values of best alpha = 0.0001 The train log loss is: 0.5278773270872483
 For values of best alpha = 0.0001 The cross validation log loss is: 1.1526650264043543
 For values of best alpha = 0.0001 The test log loss is: 1.0844189948546337

4.3.1.2. Testing the model with best hyper paramters


In [66]:

```
# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.


#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/geometric-intuition-1/
#-----
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
lr_bal_misclassified = predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y,
cv_x_onehotCoding, cv_y, clf)
```

Log loss : 1.1526650264043543
 Number of mis-classified points : 0.3966165413533835
 ----- Confusion matrix -----




----- Precision matrix (Column Sum=1) -----





----- Recall matrix (Row sum=1) -----



4.3.1.3. Feature Importance

In [67]:

```
def get_imp_feature_names(text, indices, removed_ind = []):
    word_present = 0
    tabulte_list = []
    incresingorder_ind = 0
    for i in indices:
        if i < train_gene_feature_onehotCoding.shape[1]:
            tabulte_list.append([incresingorder_ind, "Gene", "Yes"])
        elif i < 18:
            tabulte_list.append([incresingorder_ind, "Variation", "Yes"])
        if ((i > 17) & (i not in removed_ind)) :
            word = train_text_features[i]
            yes_no = True if word in text.split() else False
            if yes_no:
                word_present += 1
                tabulte_list.append([incresingorder_ind, train_text_features[i], yes_no])
            incresingorder_ind += 1
    print(word_present, "most important features are present in our query point")
    print("-"*50)
    print("The features that are most important of the ", predicted_cls[0], " class:")
    print(tabulate(tabulte_list, headers=["Index", "Feature name", "Present or Not"]))
```

4.3.1.3.1. Correctly Classified point

In [68]:

```
# from tabulate import tabulate
```

```

clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_x_onehotCoding, train_y)
test_point_index = 1
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
      np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]), 4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_) [predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index], test_df['Gene'].iloc[test_point_index], test_df['Variation']
      .iloc[test_point_index], no_feature)

```

```

Predicted Class : 7
Predicted Class Probabilities: [[0.2346 0.104  0.0378 0.1562 0.0766 0.046  0.3272 0.006  0.0115]]
Actual Class : 2

```

Out of the top 500 features 0 are present in query point

4.3.1.3.2. Incorrectly Classified point

In [69]:

```

test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
      np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]), 4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_) [predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index], test_df['Gene'].iloc[test_point_index], test_df['Variation']
      .iloc[test_point_index], no_feature)

```

```

Predicted Class : 4
Predicted Class Probabilities: [[0.3564 0.0358 0.0134 0.4356 0.0681 0.0387 0.0424 0.0042 0.0053]]
Actual Class : 4

```

209 Text feature [112] present in test data point [True]
291 Text feature [1213] present in test data point [True]
Out of the top 500 features 2 are present in query point

4.3.2. Without Class balancing

4.3.2.1. Hyper parameter tuning

In [70]:

```

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/geometric-in-tuition-1/

```

```

#-----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10 ** x for x in range(-6, 1)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

# summarizing data
lr_best_alpha = alpha[best_alpha]
lr_encoding = "one hot"

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
lr_train_log_loss = log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train,
predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
lr_cv_log_loss = log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_lo
ss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(test_x_onehotCoding)
lr_test_log_loss = log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, p
redict_y, labels=clf.classes_, eps=1e-15))

```

```


for alpha = 1e-06
Log Loss : 1.437913027879037
for alpha = 1e-05
Log Loss : 1.3397112211577944
for alpha = 0.0001
Log Loss : 1.1420789471742039
for alpha = 0.001
Log Loss : 1.1518386472331532

```

```

for alpha = 0.01
Log Loss : 1.276333430941173
for alpha = 0.1
Log Loss : 1.3820832198624688
for alpha = 1
Log Loss : 1.4138914906543354

```



For values of best alpha = 0.0001 The train log loss is: 0.5072721104825129
 For values of best alpha = 0.0001 The cross validation log loss is: 1.1420789471742039
 For values of best alpha = 0.0001 The test log loss is: 1.0765681662083113

4.3.2.2. Testing model with best hyper parameters

In [71]:

```

# read more about SGDClassifier() at http://scikit-
learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_i
ter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0
=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)


# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.


#-----
# video link:
#-----

clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
lr_misclassified = predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y,
cv_x_onehotCoding, cv_y, clf)


```

Log loss : 1.1420789471742039
 Number of mis-classified points : 0.41353383458646614
 ----- Confusion matrix -----






----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



4.3.2.3. Feature Importance, Correctly Classified point

In [72]:

```
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_x_onehotCoding,train_y)
test_point_index = 1
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,:no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index],test_df['Gene'].iloc[test_point_index],test_df['Variation']
.iloc[test_point_index], no_feature)
```

```
Predicted Class : 7
Predicted Class Probabilities: [[0.2428 0.1027 0.0329 0.1323 0.0759 0.0495 0.3372 0.0132 0.0136]]
Actual Class : 2
-----
Out of the top 500 features 0 are present in query point
```

4.3.2.4. Feature Importance, Inorrectly Classified point

In [73]:

```
test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
      np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index],test_df['Gene'].iloc[test_point_index],test_df['Variation']
.iloc[test_point_index], no_feature)
```

```
Predicted Class : 4
Predicted Class Probabilities: [[0.3524 0.0371 0.01    0.4417 0.0669 0.0377 0.0477 0.0033 0.0031]]
Actual Class : 4
-----
168 Text feature [1213] present in test data point [True]
233 Text feature [112] present in test data point [True]
Out of the top 500 features 2 are present in query point
```

4.4. Linear Support Vector Machines

4.4.1. Hyper paramter tuning

In [74]:

```
# read more about support vector machines with linear kernals here http://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html

# -----
# default parameters
# SVC(C=1.0, kernel='rbf', degree=3, gamma='auto', coef0=0.0, shrinking=True, probability=False, tol=0.001,
# cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', random_state=None)

# Some of methods of SVM()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/mathematical-derivation-copy-8/
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
```

```
#-----
alpha = [10 ** x for x in range(-5, 3)]
cv_log_error_array = []
for i in alpha:
    print("for C =", i)
    # clf = SVC(C=i, kernel='linear', probability=True, class_weight='balanced')
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='hinge', random_state=42)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
# clf = SVC(C=i, kernel='linear', probability=True, class_weight='balanced')
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='hinge', random_state=42)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)


# summarizing data
svm_best_alpha = alpha[best_alpha]
svm_encoding = "one hot"


predict_y = sig_clf.predict_proba(train_x_onehotCoding)
svm_train_log_loss = log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
svm_cv_log_loss = log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(test_x_onehotCoding)
svm_test_log_loss = log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
```

```
for C = 1e-05
Log Loss : 1.4071087671846414
for C = 0.0001
Log Loss : 1.300007591967706
for C = 0.001
Log Loss : 1.2998436621387306
for C = 0.01
Log Loss : 1.4415096810455283
for C = 0.1
Log Loss : 1.4373145068891364
for C = 1
Log Loss : 1.4385405956765203
for C = 10
Log Loss : 1.4385405978485313
for C = 100
Log Loss : 1.4385405866502794
```





For values of best alpha = 0.001 The train log loss is: 0.7013553469227897
 For values of best alpha = 0.001 The cross validation log loss is: 1.2998436621387306
 For values of best alpha = 0.001 The test log loss is: 1.24043736077651

4.4.2. Testing model with best hyper parameters

In [75]:

```
# read more about support vector machines with linear kernels here http://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html

# -----
# default parameters
# SVC(C=1.0, kernel='rbf', degree=3, gamma='auto', coef0=0.0, shrinking=True, probability=False, tol=0.001,
# cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', random_state=None)


# Some of methods of SVM()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# -----
# video link: https://www.appliedaicomse.com/course/applied-ai-course-online/lessons/mathematical-derivation-copy-8/
# -----

# clf = SVC(C=alpha[best_alpha],kernel='linear',probability=True, class_weight='balanced')
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='hinge',
random_state=42,class_weight='balanced')
svm_misclassified = predict_and_plot_confusion_matrix(train_x_onehotCoding,
train_y,cv_x_onehotCoding,cv_y, clf)
```


Log loss : 1.2998436621387306

Number of mis-classified points : 0.40601503759398494


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



4.3.3. Feature Importance

4.3.3.1. For Correctly classified point

In [76]:

```
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='hinge', random_state=42)
clf.fit(train_x_onehotCoding,train_y)
test_point_index = 1
# test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index],test_df['Gene'].iloc[test_point_index],test_df['Variation']
.iloc[test_point_index], no_feature)
```

Predicted Class : 7

Predicted Class Probabilities: [[0.204 0.0943 0.0224 0.1849 0.0626 0.0497 0.371 0.0054 0.0058]]

Actual Class : 2

492 Text feature [000] present in test data point [True]
Out of the top 500 features 1 are present in query point

4.3.3.2. For Incorrectly classified point

In [77]:

```
test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
      np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index],test_df['Gene'].iloc[test_point_index],test_df['Variation']
.iloc[test_point_index], no_feature)
```

Predicted Class : 4

Predicted Class Probabilities: [[0.2649 0.1041 0.0224 0.3109 0.0706 0.0637 0.1523 0.0059 0.0052]]

Actual Class : 4

40 Text feature [112] present in test data point [True]
471 Text feature [1213] present in test data point [True]
Out of the top 500 features 2 are present in query point

4.5 Random Forest Classifier

4.5.1. Hyper paramter tuning (With One hot Encoding)

In [78]:

```
# -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', max_depth=None, min_s
amples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None, min
impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None,
verbose=0, warm_start=False,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# predict_proba (X) Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/random-forest-and-their-construction-2/
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
# -----
```

```

# video link:
#-----

alpha = [100,200,500,1000,2000]
max_depth = [5, 10]
cv_log_error_array = []
for i in alpha:
    for j in max_depth:
        print("for n_estimators =", i,"and max depth = ", j)
        clf = RandomForestClassifier(n_estimators=i, criterion='gini', max_depth=j, random_state=42
, n_jobs=-1)
        clf.fit(train_x_onehotCoding, train_y)
        sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
        sig_clf.fit(train_x_onehotCoding, train_y)
        sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
        cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
        print("Log Loss :",log_loss(cv_y, sig_clf_probs))

'''fig, ax = plt.subplots()
features = np.dot(np.array(alpha)[: ,None],np.array(max_depth)[None]).ravel()
ax.plot(features, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[int(i/2)],max_depth[int(i%2)],str(txt)),
(features[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
'''

best_alpha = np.argmin(cv_log_error_array)
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion='gini', max_depth=max
_depth[int(best_alpha%2)], random_state=42, n_jobs=-1)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

# summarizing data
rf_1_best_alpha = best_alpha
rf_1_encoding = "one hot"

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
rf_1_train_log_loss = log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The train log loss
is:",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
rf_1_cv_log_loss = log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The cross validation log loss
is:",log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(test_x_onehotCoding)
rf_1_test_log_loss = log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The test log loss
is:",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

```

for n_estimators = 100 and max depth = 5
Log Loss : 1.3475874702229613
for n_estimators = 100 and max depth = 10
Log Loss : 1.3278965146506423
for n_estimators = 200 and max depth = 5
Log Loss : 1.3316136310087159
for n_estimators = 200 and max depth = 10
Log Loss : 1.319723891280248
for n_estimators = 500 and max depth = 5
Log Loss : 1.3197247698040182
for n_estimators = 500 and max depth = 10
Log Loss : 1.3133392096988323
for n_estimators = 1000 and max depth = 5
Log Loss : 1.3178367460504865
for n_estimators = 1000 and max depth = 10
Log Loss : 1.313050972390351
for n_estimators = 2000 and max depth = 5
Log Loss : 1.3160513445393514
for n_estimators = 2000 and max depth = 10

```

Log Loss : 1.311889243831091

For values of best estimator = 2000 The train log loss is: 1.0205284564046668

For values of best estimator = 2000 The cross validation log loss is: 1.311889243831091

For values of best estimator = 2000 The test log loss is: 1.2567561360308785

4.5.2. Testing model with best hyper parameters (One Hot Encoding)

In [79]:

```
# -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', max_depth=None, min_
samples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None, min_
impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None,
verbose=0, warm_start=False,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# predict_proba (X) Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).


# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/random-fore
t-and-their-construction-2/
# -----

clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion='gini', max_depth=max_
depth[int(best_alpha*2)], random_state=42, n_jobs=-1)
rf_1_misclassified = predict_and_plot_confusion_matrix(train_x_onehotCoding,
train_y,cv_x_onehotCoding,cv_y, clf)
```


Log loss : 1.311889243831091


Number of mis-classified points : 0.4718045112781955

----- Confusion matrix -----




----- Precision matrix (Column Sum=1) -----





----- Recall matrix (Row sum=1) -----



4.5.3. Feature Importance

4.5.3.1. Correctly Classified point

In [80]:

```
# test_point_index = 10
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion='gini', max_depth=max_depth[int(best_alpha*2)], random_state=42, n_jobs=-1)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

test_point_index = 1
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
      np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("-"*50)
get_impfeature_names(indices[:no_feature], test_df['TEXT'].iloc[test_point_index], test_df['Gene'].iloc[test_point_index], test_df['Variation'].iloc[test_point_index], no_feature)
```

Predicted Class : 7

Predicted Class Probabilities: [[0.1907 0.1214 0.0246 0.1941 0.0649 0.0684 0.3222 0.0064 0.0074]]

Actual Class : 2

Out of the top 100 features 0 are present in query point

4.5.3.2. Incorrectly Classified point

In [81]:

```
test_point_index = 100
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("-"*50)
get_impfeature_names(indices[:no_feature], test_df['TEXT'].iloc[test_point_index],test_df['Gene'].
iloc[test_point_index],test_df['Variation'].iloc[test_point_index], no_feature)
```

Predicted Class : 1

Predicted Class Probabilities: [[0.3313 0.0766 0.0184 0.3183 0.0604 0.0554 0.1282 0.0053 0.006]]

Actual Class : 4

Out of the top 100 features 0 are present in query point

4.5.3. Hyper paramter tuning (With Response Coding)

In [82]:

```
# -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', max_depth=None, min_s
amples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None, min
impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None,
verbose=0, warm_start=False,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# predict_proba (X) Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/random-forest-and-their-construction-2/
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10,50,100,200,500,1000]
max_depth = [2,3,5,10]
cv_log_error_array = []
for i in alpha:
    for j in max_depth:
        print("for n_estimators =", i,"and max depth = ", j)
        clf = RandomForestClassifier(n_estimators=i, criterion='gini', max_depth=j, random_state=42
, n_jobs=-1)
        clf.fit(train_x_responseCoding, train_y)
```

```

sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)
sig_clf_probs = sig_clf.predict_proba(cv_x_responseCoding)
cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
print("Log Loss :", log_loss(cv_y, sig_clf_probs))
'''
fig, ax = plt.subplots()
features = np.dot(np.array(alpha[:, None], np.array(max_depth)[None]).ravel(),
ax.plot(features, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[int(i/4)], max_depth[int(i%4)], str(txt)),
(features[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
'''

best_alpha = np.argmin(cv_log_error_array)
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/4)], criterion='gini', max_depth=max
_depth[int(best_alpha/4)], random_state=42, n_jobs=-1)
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

# summarizing data
rf_best_alpha = None
rf_encoding = "Response"

predict_y = sig_clf.predict_proba(train_x_responseCoding)
rf_train_log_loss = log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The train log loss is:", log_loss(y
_train, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(cv_x_responseCoding)
rf_cv_log_loss = log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The cross validation log loss is:"
, log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(test_x_responseCoding)
rf_test_log_loss = log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The test log loss is:", log_loss(y
_test, predict_y, labels=clf.classes_, eps=1e-15))

```

```

for n_estimators = 10 and max depth = 2
Log Loss : 2.273544813083507
for n_estimators = 10 and max depth = 3
Log Loss : 1.7884768301624645
for n_estimators = 10 and max depth = 5
Log Loss : 1.5293694431872824
for n_estimators = 10 and max depth = 10
Log Loss : 1.8158170664765045
for n_estimators = 50 and max depth = 2
Log Loss : 1.8142836492644847
for n_estimators = 50 and max depth = 3
Log Loss : 1.5103739217733743
for n_estimators = 50 and max depth = 5
Log Loss : 1.4709509481795053
for n_estimators = 50 and max depth = 10
Log Loss : 1.8145770053270156
for n_estimators = 100 and max depth = 2
Log Loss : 1.6154321026576617
for n_estimators = 100 and max depth = 3
Log Loss : 1.5748970214382243
for n_estimators = 100 and max depth = 5
Log Loss : 1.3519182392743645
for n_estimators = 100 and max depth = 10
Log Loss : 1.81363920078842
for n_estimators = 200 and max depth = 2
Log Loss : 1.663321158212204
for n_estimators = 200 and max depth = 3
Log Loss : 1.5707324468249035
for n_estimators = 200 and max depth = 5
Log Loss : 1.4026579207504937
for n_estimators = 200 and max depth = 10

```

```

for n_estimators = 200 and max depth = 10
Log Loss : 1.8502392723388763
for n_estimators = 500 and max depth = 2
Log Loss : 1.756639773364519
for n_estimators = 500 and max depth = 3
Log Loss : 1.645109085600542
for n_estimators = 500 and max depth = 5
Log Loss : 1.4323133967303667
for n_estimators = 500 and max depth = 10
Log Loss : 1.8849344494978828
for n_estimators = 1000 and max depth = 2
Log Loss : 1.715936199027195
for n_estimators = 1000 and max depth = 3
Log Loss : 1.6488750463651782
for n_estimators = 1000 and max depth = 5
Log Loss : 1.406735070147225
for n_estimators = 1000 and max depth = 10
Log Loss : 1.8734321784764465
For values of best alpha = 100 The train log loss is: 0.05143647638254964
For values of best alpha = 100 The cross validation log loss is: 1.3519182392743645
For values of best alpha = 100 The test log loss is: 1.2745618371023875

```

4.5.4. Testing model with best hyper parameters (Response Coding)

In [83]:

```

# -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', max_depth=None, min_
samples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None, min_
impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None,
verbose=0, warm_start=False,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# predict_proba (X) Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/random-forest-and-their-construction-2/
# -----


clf = RandomForestClassifier(max_depth=max_depth[int(best_alpha%4)],
n_estimators=alpha[int(best_alpha/4)], criterion='gini', max_features='auto', random_state=42)
rf_misclassified = predict_and_plot_confusion_matrix(train_x_responseCoding,
train_y,cv_x_responseCoding,cv_y, clf)


```

```


Log loss : 1.3519182392743645
Number of mis-classified points : 0.46616541353383456
----- Confusion matrix -----

```






----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



4.5.5. Feature Importance

4.5.5.1. Correctly Classified point

In [84]:

```
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/4)], criterion='gini', max_depth=max_depth[int(best_alpha*4)], random_state=42, n_jobs=-1)
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

test_point_index = 1
no_feature = 27
predicted_cls = sig_clf.predict(test_x_responseCoding[test_point_index].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(test_x_responseCoding[test_point_index].reshape(1,-1)).4))
```

```

np.round(sig_clf.predict_proba(test_x_responseCoding[test_point_index].reshape(1,-1)),4)
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("--*50)
for i in indices:
    if i<9:
        print("Gene is important feature")
    elif i<18:
        print("Variation is important feature")
    else:
        print("Text is important feature")

```

Predicted Class : 8

Predicted Class Probabilities: [[0.3023 0.016 0.0267 0.0451 0.0217 0.0437 0.0107 0.4188 0.1151]]

Actual Class : 2

```

-----
Variation is important feature
Variation is important feature
Variation is important feature
Variation is important feature
Gene is important feature
Variation is important feature
Variation is important feature
Text is important feature
Text is important feature
Gene is important feature
Text is important feature
Text is important feature
Text is important feature
Text is important feature
Gene is important feature
Gene is important feature
Variation is important feature
Text is important feature
Gene is important feature
Variation is important feature
Gene is important feature
Text is important feature
Gene is important feature
Gene is important feature
Variation is important feature
Text is important feature
Text is important feature
Gene is important feature

```

4.5.5.2. Incorrectly Classified point

In [85]:

```

test_point_index = 100
predicted_cls = sig_clf.predict(test_x_responseCoding[test_point_index].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(test_x_responseCoding[test_point_index].reshape(1,-1)),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("--*50)
for i in indices:
    if i<9:
        print("Gene is important feature")
    elif i<18:
        print("Variation is important feature")
    else:
        print("Text is important feature")

```

Predicted Class : 4

Predicted Class Probabilities: [[0.2949 0.0159 0.1269 0.3937 0.0449 0.0664 0.0085 0.0276 0.0212]]

Actual Class : 4

```

-----
Variation is important feature
Variation is important feature
Variation is important feature
Variation is important feature
Gene is important feature
Variation is important feature

```

Variation is important feature
Text is important feature
Text is important feature
Gene is important feature
Text is important feature
Text is important feature
Text is important feature
Gene is important feature
Gene is important feature
Variation is important feature
Text is important feature
Gene is important feature
Variation is important feature
Gene is important feature
Text is important feature
Gene is important feature
Gene is important feature
Variation is important feature
Text is important feature
Text is important feature
Gene is important feature

4.7 Stack the models

4.7.1 testing with hyper parameter tuning

In [86]:

```
# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/geometric-intuition-1/
#-----

# read more about support vector machines with linear kernels here http://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html
# -----
# default parameters
# SVC(C=1.0, kernel='rbf', degree=3, gamma='auto', coef0=0.0, shrinking=True, probability=False, tol=0.001,
# cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', random_state=None)

# Some of methods of SVM()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/mathematical-derivation-copy-8/
# -----

# read more about support vector machines with linear kernels here http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html
# -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', max_depth=None, min_samples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None, min_
```

```

impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None,
verbose=0, warm_start=False,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# predict_proba (X) Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/random-forest-and-their-construction-2/
# -----

clf1 = SGDClassifier(alpha=0.001, penalty='l2', loss='log', class_weight='balanced', random_state=0)
clf1.fit(train_x_onehotCoding, train_y)
sig_clf1 = CalibratedClassifierCV(clf1, method="sigmoid")

clf2 = SGDClassifier(alpha=1, penalty='l2', loss='hinge', class_weight='balanced', random_state=0)
clf2.fit(train_x_onehotCoding, train_y)
sig_clf2 = CalibratedClassifierCV(clf2, method="sigmoid")

clf3 = MultinomialNB(alpha=0.001)
clf3.fit(train_x_onehotCoding, train_y)
sig_clf3 = CalibratedClassifierCV(clf3, method="sigmoid")

sig_clf1.fit(train_x_onehotCoding, train_y)
LR = (log_loss(cv_y, sig_clf1.predict_proba(cv_x_onehotCoding)))
print("Logistic Regression : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf1.predict_proba(cv_x_onehotCoding))))

sig_clf2.fit(train_x_onehotCoding, train_y)
SVM = (log_loss(cv_y, sig_clf2.predict_proba(cv_x_onehotCoding)))
print("Support vector machines : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf2.predict_proba(cv_x_onehotCoding))))

sig_clf3.fit(train_x_onehotCoding, train_y)
NB = (log_loss(cv_y, sig_clf3.predict_proba(cv_x_onehotCoding)))
print("Naive Bayes : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf3.predict_proba(cv_x_onehotCoding))))

print("-"*50)
alpha = [0.0001,0.001,0.01,0.1,1,10]
best_alpha_loss = 999
for i in alpha:
    lr = LogisticRegression(C=i)
    sclf = StackingClassifier(classifiers=[sig_clf1, sig_clf2, sig_clf3], meta_classifier=lr, use_probas=True)
    sclf.fit(train_x_onehotCoding, train_y)
    print("Stacking Classifier : for the value of alpha: %f Log Loss: %0.3f" % (i, log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))))
    log_error = log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))
    if best_alpha_loss > log_error:
        best_alpha_loss = log_error
        best_alpha = i

```

```

Logistic Regression : Log Loss: 1.17
Support vector machines : Log Loss: 1.44
Naive Bayes : Log Loss: 1.30

```

```

-----
Stacking Classifier : for the value of alpha: 0.000100 Log Loss: 2.179
Stacking Classifier : for the value of alpha: 0.001000 Log Loss: 2.045
Stacking Classifier : for the value of alpha: 0.010000 Log Loss: 1.571
Stacking Classifier : for the value of alpha: 0.100000 Log Loss: 1.217
Stacking Classifier : for the value of alpha: 1.000000 Log Loss: 1.347
Stacking Classifier : for the value of alpha: 10.000000 Log Loss: 1.772

```



```
stack_best_alpha = best_alpha
stack_encoding = "one hot"
print(best_alpha_loss)
```

4.7.2 testing the model with the best hyper parameters

```
lr = LogisticRegression(C=0.1)
sclf = StackingClassifier(classifiers=[sig_clf1, sig_clf2, sig_clf3], meta_classifier=lr, use_proba=True)
sclf.fit(train_x_onehotCoding, train_y)



log_error = log_loss(train_y, sclf.predict_proba(train_x_onehotCoding))
stack_train_log_loss = log_error
print("Log loss (train) on the stacking classifier :", log_error)


log_error = log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))
stack_cv_log_loss = log_error
print("Log loss (CV) on the stacking classifier :", log_error)

log_error = log_loss(test_y, sclf.predict_proba(test_x_onehotCoding))
stack_test_log_loss = log_error
print("Log loss (test) on the stacking classifier :", log_error)


stack_misclassified = np.count_nonzero((sclf.predict(test_x_onehotCoding) - test_y)) / test_y.shape[0]
print("Number of misclassified point :", stack_misclassified)
plot_confusion_matrix(test_y=test_y, predict_y=sclf.predict(test_x_onehotCoding))
```

```
----- Confusion matrix -----
```





```
----- Recall matrix (Row sum=1) -----
```



4.7.3 Maximum Voting classifier

In [89]:

```
#Refer:http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.VotingClassifier.html
from sklearn.ensemble import VotingClassifier
vclf = VotingClassifier(estimators=[('lr', sig_clf1), ('svc', sig_clf2), ('rf', sig_clf3)], voting=
'soft')
vclf.fit(train_x_onehotCoding, train_y)

voting_best_alpha = None
voting_encoding = "One Hot"
voting_train_log_loss = log_loss(train_y, vclf.predict_proba(train_x_onehotCoding))
print("Log loss (train) on the VotingClassifier :", log_loss(train_y,
vclf.predict_proba(train_x_onehotCoding)))


voting_cv_log_loss = log_loss(cv_y, vclf.predict_proba(cv_x_onehotCoding))
print("Log loss (CV) on the VotingClassifier :", log_loss(cv_y,
vclf.predict_proba(cv_x_onehotCoding)))


voting_test_log_loss = log_loss(test_y, vclf.predict_proba(test_x_onehotCoding))
print("Log loss (test) on the VotingClassifier :", log_loss(test_y,
vclf.predict_proba(test_x_onehotCoding)))

voting_misclassified = np.count_nonzero((vclf.predict(test_x_onehotCoding)- test_y))/test_y.shape[0
]
print("Number of missclassified point :", voting_misclassified)
plot_confusion_matrix(test_y=test_y, predict_y=vclf.predict(test_x_onehotCoding))
```


```
Log loss (train) on the VotingClassifier : 0.8144237343471504
Log loss (CV) on the VotingClassifier : 1.2519537206391715
Log loss (test) on the VotingClassifier : 1.2038507556613278
Number of missclassified point : 0.3819548872180451
```

```
----- Confusion matrix -----
```





Precision matrix (Column Sum=1)



Recall matrix (Row sum=1)



Task2 Summary

In [90]:

```
from prettytable import PrettyTable

# to reference from http://zetcode.com/python/prettytable/
```

```
summary = PrettyTable()
```

```
In [91]:
```

```
summary = PrettyTable()
summary.field_names = ["Model", "Encoding", "Best Alpha", "Train logloss", "CV logloss", "Test logloss", "MisClassified"]
```

```
In [92]:
```

```
summary.add_row(["Naive Bayes", nb_encoding, nb_best_alpha, round(nb_train_log_loss, 3), round(nb_cv_log_loss, 3), round(nb_test_log_loss, 3), nb_misclassified])

summary.add_row(["Logistic R.", lr_encoding, lr_best_alpha, round(lr_train_log_loss, 3), round(lr_cv_log_loss, 3), round(lr_test_log_loss, 3), lr_misclassified])
summary.add_row(["Logistic Balanced", lr_bal_encoding, lr_bal_best_alpha, round(lr_bal_train_log_loss, 3), round(lr_bal_cv_log_loss, 3), round(lr_bal_test_log_loss, 3), lr_bal_misclassified])
summary.add_row(["Linear SVM", svm_encoding, svm_best_alpha, round(svm_train_log_loss, 3), round(svm_cv_log_loss, 3), round(svm_test_log_loss, 3), svm_misclassified])

summary.add_row(["KNN classifier", knn_encoding, knn_best_alpha, round(knn_train_log_loss, 3), round(knn_cv_log_loss, 3), round(knn_test_log_loss, 3), knn_misclassified])

summary.add_row(["Random Forest", rf_1_encoding, rf_1_best_alpha, round(rf_1_train_log_loss, 3), round(rf_1_cv_log_loss, 3), round(rf_1_test_log_loss, 3), rf_1_misclassified])
summary.add_row(["Random Forest", rf_encoding, rf_best_alpha, round(rf_train_log_loss, 3), round(rf_cv_log_loss, 3), round(rf_test_log_loss, 3), rf_misclassified])
summary.add_row(["Stacking", stack_encoding, stack_best_alpha, round(stack_train_log_loss, 3), round(stack_cv_log_loss, 3), round(stack_test_log_loss, 3), stack_misclassified])
summary.add_row(["Max. Voting", voting_encoding, voting_best_alpha, round(voting_train_log_loss, 3), round(voting_cv_log_loss, 3), round(voting_test_log_loss, 3), voting_misclassified])
```

```
In [93]:
```

```
print("Model and their performance...\n")
print(summary)
```

Model and their performance...

```
+-----+-----+-----+-----+-----+-----+-----+
|      Model      | Encoding | Best Alpha | Train logloss | CV logloss | Test logloss | MisClassified |
+-----+-----+-----+-----+-----+-----+-----+
| Naive Bayes     | One hot  | 1          | 0.778         | 1.209      | 1.163        | 0.439849240601504 |
| Logistic R.     | one hot  | 0.0001     | 0.507         | 1.142      | 1.077        | 0.4135338458646614 |
| Logistic Balanced | One hot  | 0.0001     | 0.528         | 1.153      | 1.084        | 0.3966165413533835 |
| Linear SVM      | one hot  | 0.001      | 0.701         | 1.3         | 1.24         | 0.4060150759398494 |
| KNN classifier  | Response | 5          | 0.487         | 1.045      | 1.031        | 0.3515037939849626 |
| Random Forest   | one hot  | 9          | 1.021         | 1.312      | 1.257        | 0.471804112781955 |
| Random Forest   | Response | None       | 0.051         | 1.352      | 1.275        | 0.4661654353383456 |
| Stacking        | one hot  | 0.1        | 0.572         | 1.217      | 1.191        | 0.384962060150376 |
| Max. Voting     | One Hot  | None       | 0.814         | 1.252      | 1.204        | 0.381954872180451 |
+-----+-----+-----+-----+-----+-----+-----+
|      Model      | Encoding | Best Alpha | Train logloss | CV logloss | Test logloss | MisClassified |
+-----+-----+-----+-----+-----+-----+-----+
```

Task3 Logistic regression with CountVectorizer

Gene Feature

In [94]:

```
# one-hot encoding of Gene feature.
gene_vectorizer = CountVectorizer()
train_gene_feature_onehotCoding = gene_vectorizer.fit_transform(train_df['Gene'])
test_gene_feature_onehotCoding = gene_vectorizer.transform(test_df['Gene'])
cv_gene_feature_onehotCoding = gene_vectorizer.transform(cv_df['Gene'])
```

```
# one-hot encoding of Gene feature.
gene_vectorizer = CountVectorizer()
train_gene_feature_onehotCoding = gene_vectorizer.fit_transform(train_df['Gene'])
test_gene_feature_onehotCoding = gene_vectorizer.transform(test_df['Gene'])
cv_gene_feature_onehotCoding = gene_vectorizer.transform(cv_df['Gene'])
```

Variation

In [95]:

```
# alpha is used for laplace smoothing
alpha = 1
# train gene feature
train_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", train_df))
# test gene feature
test_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", test_df))
# cross validation gene feature
cv_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", cv_df))

# one-hot encoding of variation feature.
variation_vectorizer = CountVectorizer()
train_variation_feature_onehotCoding = variation_vectorizer.fit_transform(train_df['Variation'])
test_variation_feature_onehotCoding = variation_vectorizer.transform(test_df['Variation'])
cv_variation_feature_onehotCoding = variation_vectorizer.transform(cv_df['Variation'])
```

Text Feature

In [96]:

```
# building a CountVectorizer with all the words that occurred minimum 3 times in train data
text_vectorizer = CountVectorizer(min_df=3, ngram_range=(1, 2))
train_text_feature_onehotCoding = text_vectorizer.fit_transform(train_df['TEXT'])

# getting all the feature names (words)
train_text_features = text_vectorizer.get_feature_names()

# train_text_feature_onehotCoding.sum(axis=0).A1 will sum every row and returns (1*number of features) vector
train_text_fea_counts = train_text_feature_onehotCoding.sum(axis=0).A1

# zip(list(text_features), text_fea_counts) will zip a word with its number of times it occurred
text_fea_dict = dict(zip(list(train_text_features), train_text_fea_counts))

# response coding of text features
train_text_feature_responseCoding = get_text_responsecoding(train_df)
test_text_feature_responseCoding = get_text_responsecoding(test_df)
cv_text_feature_responseCoding = get_text_responsecoding(cv_df)

# https://stackoverflow.com/a/16202486
# we convert each row values such that they sum to 1
train_text_feature_responseCoding = (train_text_feature_responseCoding.T / train_text_feature_responseCoding.sum(axis=1)).T
test_text_feature_responseCoding = (test_text_feature_responseCoding.T / test_text_feature_responseCoding.sum(axis=1)).T
cv_text_feature_responseCoding = (cv_text_feature_responseCoding.T / cv_text_feature_responseCoding.sum(axis=1)).T
```

In [97]:

```
# don't forget to normalize every feature
train_text_feature_onehotCoding = normalize(train_text_feature_onehotCoding, axis=0)

# we use the same vectorizer that was trained on train data
test_text_feature_onehotCoding = text_vectorizer.transform(test_df['TEXT'])
# don't forget to normalize every feature
test_text_feature_onehotCoding = normalize(test_text_feature_onehotCoding, axis=0)

# we use the same vectorizer that was trained on train data
cv_text_feature_onehotCoding = text_vectorizer.transform(cv_df['TEXT'])
# don't forget to normalize every feature
cv_text_feature_onehotCoding = normalize(cv_text_feature_onehotCoding, axis=0)
```

Stacking the Features

In [98]:

```
train_gene_var_onehotCoding =
hstack((train_gene_feature_onehotCoding, train_variation_feature_onehotCoding))
test_gene_var_onehotCoding =
hstack((test_gene_feature_onehotCoding, test_variation_feature_onehotCoding))
cv_gene_var_onehotCoding = hstack((cv_gene_feature_onehotCoding, cv_variation_feature_onehotCoding))

train_x_onehotCoding = hstack((train_gene_var_onehotCoding, train_text_feature_onehotCoding)).tocsr()
train_y = np.array(list(train_df['Class']))

test_x_onehotCoding = hstack((test_gene_var_onehotCoding, test_text_feature_onehotCoding)).tocsr()
test_y = np.array(list(test_df['Class']))

cv_x_onehotCoding = hstack((cv_gene_var_onehotCoding, cv_text_feature_onehotCoding)).tocsr()
cv_y = np.array(list(cv_df['Class']))

train_gene_var_responseCoding =
np.hstack((train_gene_feature_responseCoding, train_variation_feature_responseCoding))
test_gene_var_responseCoding =
np.hstack((test_gene_feature_responseCoding, test_variation_feature_responseCoding))
cv_gene_var_responseCoding =
np.hstack((cv_gene_feature_responseCoding, cv_variation_feature_responseCoding))

train_x_responseCoding = np.hstack((train_gene_var_responseCoding,
train_text_feature_responseCoding))
test_x_responseCoding = np.hstack((test_gene_var_responseCoding, test_text_feature_responseCoding))
cv_x_responseCoding = np.hstack((cv_gene_var_responseCoding, cv_text_feature_responseCoding))
```

Print

In [99]:

```
print("One hot encoding features :")
print("(number of data points * number of features) in train data = ", train_x_onehotCoding.shape)
print("(number of data points * number of features) in test data = ", test_x_onehotCoding.shape)
print("(number of data points * number of features) in cross validation data =", cv_x_onehotCoding.shape)
```

```
One hot encoding features :
(number of data points * number of features) in train data = (2124, 786163)
(number of data points * number of features) in test data = (665, 786163)
(number of data points * number of features) in cross validation data = (532, 786163)
```

In [100]:

```
print(" Response encoding features :")
print("(number of data points * number of features) in train data = ", train_x_responseCoding.shape)
print("(number of data points * number of features) in test data = ", test_x_responseCoding.shape)
```

```
print("(number of data points * number of features) in cross validation data =",
cv_x_responseCoding.shape)
```

Response encoding features :

(number of data points * number of features) in train data = (2124, 27)

(number of data points * number of features) in test data = (665, 27)

(number of data points * number of features) in cross validation data = (532, 27)

Logistic Regression With Class Balancing

In [101]:

```
alpha = [10 ** x for x in range(-6, 3)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf.probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(y_cv, sig_clf.probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability estimates
    print("Log Loss :", log_loss(y_cv, sig_clf.probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

# summarizing data
lr_bal_best_alpha = alpha[best_alpha]
lr_bal_encoding = "One hot"

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
lr_bal_train_log_loss = log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train,
predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
lr_bal_cv_log_loss = log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv,
predict_y, labels=clf.classes_, eps=1e-15))


predict_y = sig_clf.predict_proba(test_x_onehotCoding)
lr_bal_test_log_loss = log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test,
predict_y, labels=clf.classes_, eps=1e-15))
```

```
for alpha = 1e-06
Log Loss : 1.590953902140178
for alpha = 1e-05
Log Loss : 1.5895829889498692
for alpha = 0.0001
Log Loss : 1.5931449739874912
for alpha = 0.001
Log Loss : 1.4965409060759935
for alpha = 0.01
```

```

Log Loss : 1.2689622912529397
for alpha = 0.1
Log Loss : 1.3121901465706813
for alpha = 1
Log Loss : 1.3614135580723945
for alpha = 10
Log Loss : 1.3998561917107513
for alpha = 100
Log Loss : 1.4108538002886344

```



```

For values of best alpha = 0.01 The train log loss is: 0.8263281520110783
For values of best alpha = 0.01 The cross validation log loss is: 1.2689622912529397
For values of best alpha = 0.01 The test log loss is: 1.1866852878328527

```

In [102]:

```


clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
lr_bal_misclassified = predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y, cv_x_onehotCoding, cv_y, clf)

```

```

Log loss : 1.2689622912529397
Number of mis-classified points : 0.41541353383458646
----- Confusion matrix -----


```




```


----- Precision matrix (Column Sum=1) -----

```





Recall matrix (Row sum=1)



In [103]:

```
from prettytable import PrettyTable

# to reference from http://zetcode.com/python/prettytable/
summary = PrettyTable()

summary = PrettyTable()
summary.field_names = ["Model", "Encoding", "Best Alpha", "Train logloss", "CV logloss", "Test logloss", "MisClassified"]

summary.add_row(["Logistic Balanced", lr_bal_encoding, lr_bal_best_alpha, round(lr_bal_train_log_loss, 3), round(lr_bal_cv_log_loss, 3), round(lr_bal_test_log_loss, 3), lr_bal_misclassified])
```

In [104]:

```
print(summary)
```

Model	Encoding	Best Alpha	Train logloss	CV logloss	Test logloss	MisClassified
Logistic Balanced	One hot	0.01	0.826	1.269	1.187	0.41541353383458646

Observation

1. Train Log los improved as compared to previous one.
2. No major change in model performance found.

Task4 Feature Engineering

Feature Eng. 1

1. From above we can say that gene is the most important feature
2. Let's do some feature engineering to Gene Feature
3. Since LR Balanced performed best in TASK1 we will use that one
4. Considering **Response encoding** and taking out **argmax of response encoded vector for gene feature**

In [138]:

```
# alpha is used for laplace smoothing
alpha = 1
train_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", train_df))
test_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", test_df))
cv_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", cv_df))
```

In [139]:

```
# alpha is used for laplace smoothing
alpha = 1
train_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", train_df))
test_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", test_df))
cv_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", cv_df))
```

In [140]:

```
gene_max = []
for i in train_gene_feature_responseCoding:
    gene_max.append(np.argmax(i))

test_gene_max = []
for i in test_gene_feature_responseCoding:
    test_gene_max.append(np.argmax(i))

cv_gene_max = []
for i in cv_gene_feature_responseCoding:
    cv_gene_max.append(np.argmax(i))
```

In [141]:

```
print(len(gene_max))
print(len(test_gene_max))
print(len(cv_gene_max))
```

2124
665
532

In [142]:

```
gene_max = np.array(gene_max)
test_gene_max = np.array(test_gene_max)
cv_gene_max = np.array(cv_gene_max)

# expanding dimetions
gene_max = np.expand_dims(gene_max,axis = 1)
test_gene_max = np.expand_dims(test_gene_max,axis = 1)
cv_gene_max = np.expand_dims(cv_gene_max,axis = 1)
```

In [143]:

```
from sklearn.preprocessing import StandardScaler
scalar = StandardScaler()
scalar.fit(gene_max)

gene_max = scalar.transform(gene_max)
test_gene_max = scalar.transform(test_gene_max)
```

```
cv_gene_max = scalar.transform(cv_gene_max)
```

In [144]:

```
train_gene_var_responseCoding =
np.hstack((train_gene_feature_responseCoding,train_variation_feature_responseCoding,gene_max))
test_gene_var_responseCoding =
np.hstack((test_gene_feature_responseCoding,test_variation_feature_responseCoding,test_gene_max))
cv_gene_var_responseCoding =
np.hstack((cv_gene_feature_responseCoding,cv_variation_feature_responseCoding,cv_gene_max))

train_x_responseCoding = np.hstack((train_gene_var_responseCoding,
train_text_feature_responseCoding))
test_x_responseCoding = np.hstack((test_gene_var_responseCoding, test_text_feature_responseCoding
))
cv_x_responseCoding = np.hstack((cv_gene_var_responseCoding, cv_text_feature_responseCoding))
```

In [145]:

```
alpha = [10 ** x for x in range(-6, 3)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log', random_state=42
)
    clf.fit(train_x_responseCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_responseCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_responseCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability estimates
    print("Log Loss :",log_loss(cv_y, sig_clf_probs))


fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],str(txt)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', ran
dom_state=42)
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_responseCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_loss(y_train,
predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_responseCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:",log_lo
ss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_responseCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_loss(y_test, p
redict_y, labels=clf.classes_, eps=1e-15))
```

```
for alpha = 1e-06
Log Loss : 1.1849804468672414
for alpha = 1e-05
Log Loss : 1.2612472188210133
for alpha = 0.0001
Log Loss : 1.1492455078328956
for alpha = 0.001
Log Loss : 1.2853491460035043
for alpha = 0.01
Log Loss : 1.4831385152309395
for alpha = 0.1
Log Loss : 1.6005476849096605
for alpha = 1
Log Loss : 1.6005476849096605
```

```
Log Loss : 1.6138882164755701
for alpha = 10
Log Loss : 1.6049206822487738
for alpha = 100
Log Loss : 1.5969033579228866
```



```
For values of best alpha = 0.0001 The train log loss is: 1.0025602897304609
For values of best alpha = 0.0001 The cross validation log loss is: 1.1492455078328956
For values of best alpha = 0.0001 The test log loss is: 1.138658509666389
```

Feature Eng 3

- from above we can say that considering response encoded vector completely deteriorated model performance badly
- Lets consider **Onehotencoding** with the same **argmax of response encoded gene feature**.

In [147]:

```
# for gene feature
gene_vectorizer = TfidfVectorizer()
train_gene_feature_onehotCoding = gene_vectorizer.fit_transform(train_df['Gene'])
test_gene_feature_onehotCoding = gene_vectorizer.transform(test_df['Gene'])
cv_gene_feature_onehotCoding = gene_vectorizer.transform(cv_df['Gene'])

# for variation feature
variation_vectorizer = TfidfVectorizer()
train_variation_feature_onehotCoding = variation_vectorizer.fit_transform(train_df['Variation'])
test_variation_feature_onehotCoding = variation_vectorizer.transform(test_df['Variation'])
cv_variation_feature_onehotCoding = variation_vectorizer.transform(cv_df['Variation'])

# for text feature
vectorizer = TfidfVectorizer()
train_text_feature_onehotCoding = vectorizer.fit_transform(train_df['TEXT'])
test_text_feature_onehotCoding = vectorizer.transform(test_df['TEXT'])
cv_text_feature_onehotCoding = vectorizer.transform(cv_df['TEXT'])
```

In [148]:

```
train_gene_var_onehotCoding =
hstack((train_gene_feature_onehotCoding, train_variation_feature_onehotCoding))
test_gene_var_onehotCoding =
hstack((test_gene_feature_onehotCoding, test_variation_feature_onehotCoding))
cv_gene_var_onehotCoding = hstack((cv_gene_feature_onehotCoding, cv_variation_feature_onehotCoding))
```

In [150]:

```
from scipy import sparse
gene_max=sparse.csr_matrix(gene_max)
test_gene_max = sparse.csr_matrix(test_gene_max)
cv_gene_max = sparse.csr_matrix(cv_gene_max)
```

In [151]:

```

train_x_onehotCoding = hstack((train_gene_var_onehotCoding,
train_text_feature_onehotCoding,gene_max)).tocsr()
train_y = np.array(list(train_df['Class']))

test_x_onehotCoding = hstack((test_gene_var_onehotCoding,
test_text_feature_onehotCoding,test_gene_max)).tocsr()
test_y = np.array(list(test_df['Class']))

cv_x_onehotCoding = hstack((cv_gene_var_onehotCoding, cv_text_feature_onehotCoding,cv_gene_max)).tocsr()
cv_y = np.array(list(cv_df['Class']))

```

In [152]:

```

alpha = [10 ** x for x in range(-6, 3)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability estimates
    print("Log Loss :",log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],str(txt)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)


predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:",log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

```

for alpha = 1e-06
Log Loss : 1.110794943085918
for alpha = 1e-05
Log Loss : 1.0699777885310346
for alpha = 0.0001
Log Loss : 1.015569262975881
for alpha = 0.001
Log Loss : 1.0765069714750855
for alpha = 0.01
Log Loss : 1.269412050233549
for alpha = 0.1
Log Loss : 1.464478652459595
for alpha = 1
Log Loss : 1.5535821142665585
for alpha = 10
Log Loss : 1.5657459167936663
for alpha = 100
Log Loss : 1.5670130235912416

```




For values of best alpha = 0.0001 The train log loss is: 0.4419919628846869
 For values of best alpha = 0.0001 The cross validation log loss is: 1.015569262975881
 For values of best alpha = 0.0001 The test log loss is: 0.980974421883897


In [153]:


```
lr_bal_misclassified_f = predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y,
cv_x_onehotCoding, cv_y, clf)
```

Log loss : 1.015569262975881
 Number of mis-classified points : 0.34774436090225563
 ----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----





In [158]:

```
from prettytable import PrettyTable

summary = PrettyTable()
summary = PrettyTable()
summary.field_names = ["FeatureEng.", "Model", "Encoding", "BestAlpha", "Trainlogloss", "CVlogloss", "Testlogloss", "MissClassified"]
summary.add_row(["Feature eng. 1", "Logistic Balanced", "response", 0.0001, 1.002, 1.14, 1.1386, 0.4582])
summary.add_row(["Feature eng. 2", "Logistic Balanced", "onehot + response", 0.0001, 0.4419, 1.015, 0.98, 0.3477])
print(summary)
```

FeatureEng.	Model	Encoding	BestAlpha	Trainlogloss	CVlogloss	Testlogloss	MissClassified
Feature eng. 1	Logistic Balanced	response	0.0001	1.002	1.14	1.1386	0.4582
Feature eng. 2	Logistic Balanced	onehot + response	0.0001	0.4419	1.015	0.98	0.3477

Observation:

- We have got **train and test log loss < 1** as asked.
- Also the **mis-classification is very less** as compare to all the models considered

CaseStudy Flow:

=====

1. The objective of the case study was to Classify the given genetic variations/mutations based on evidence from text-based clinical literature.
2. The case study demands very high interpretability and probabilistic outputs
3. Dataset contains ID, Gene, Variation, Class and Text as feature.
4. On EDA on class label it was found that **distribution of classes were not balanced**. More data pts. was present in classes 1, 2, 4 and 7 as compared to other.
5. **Gene feature found to be most important feature** followed by variation and text.
6. Features are encoded as response encode and onehot encoding.
7. Various ML models were tried and tested to obtain the best results.
8. For correctly classified pts. **difference in probabilities of classes were high** which is as expected.
9. **Random Forest with onehot encoding** took the longest to run.

10. **Random Forest with onehot encoding** showed incorrectly classified pt. as correctly classified.
11. As **LR with onehot encoding** performed best among all the models. So , it was considered during feature engineering.
12. During Feature engineering, **argmax of response encoded gene feature** was considered along with one hot encoded feature.
13. All the results are summarized in tabular format and observation are noted whenever necessary.

In []: